



# Planning and Scheduling: Hierarchical Task Network Planning



Hochschule  
Bonn-Rhein-Sieg



Prof. Dr.-Ing. Gerhard K. Kraetzschmar

# Acknowledgements



- These slides are based on those slides by Dana Nau, Gerhard Wickler, Hai Hoang, José Luis Ambite
- Several improvements were applied by Iman Awaad

- We may already have an idea how to go about solving problems in a planning domain
- E.g.: travel to a destination that's far away:
  - Domain-independent planner:
    - many combinations of vehicles and routes
  - Experienced human: small number of “recipes”, e.g. for flying:
    - buy ticket from local airport to remote airport
    - travel to local airport
    - fly to remote airport
    - travel to final destination
- How to enable planning systems to make use of such recipes?

# Two Approaches



- Control rules (Chapter 10 in the book):
  - Write rules to prune every action that does not fit the recipe
- Hierarchical Task Network (HTN) planning:
  - Describe the actions and subtasks that do fit the recipe

# Control Rules v HTN Planning

- Control rules (Chapter 10):
  - Write rules to prune every action that **does not** fit the recipe
- Hierarchical Task Network (HTN) planning:
  - Describe the actions and subtasks that **do** fit the recipe
- Objective of HTN planning: perform a given set of tasks
- Inputs include:
  - **Operators**: can directly perform a primitive task
  - **Methods**: recipes for decomposing a complex/non-primitive task into simpler non-primitive or primitive subtasks
- Planning process:
  - **Decompose** non-primitive tasks **recursively** until primitive tasks are reached

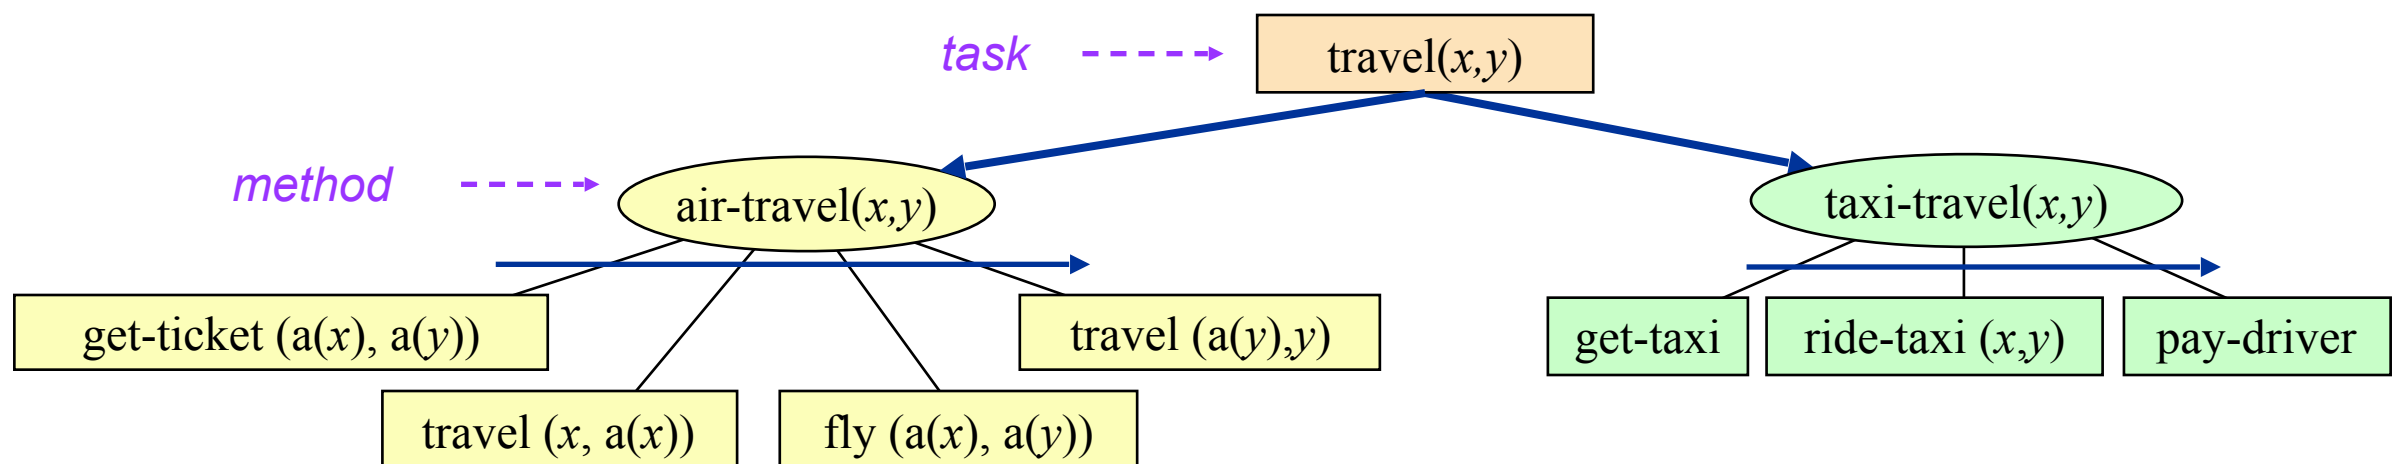
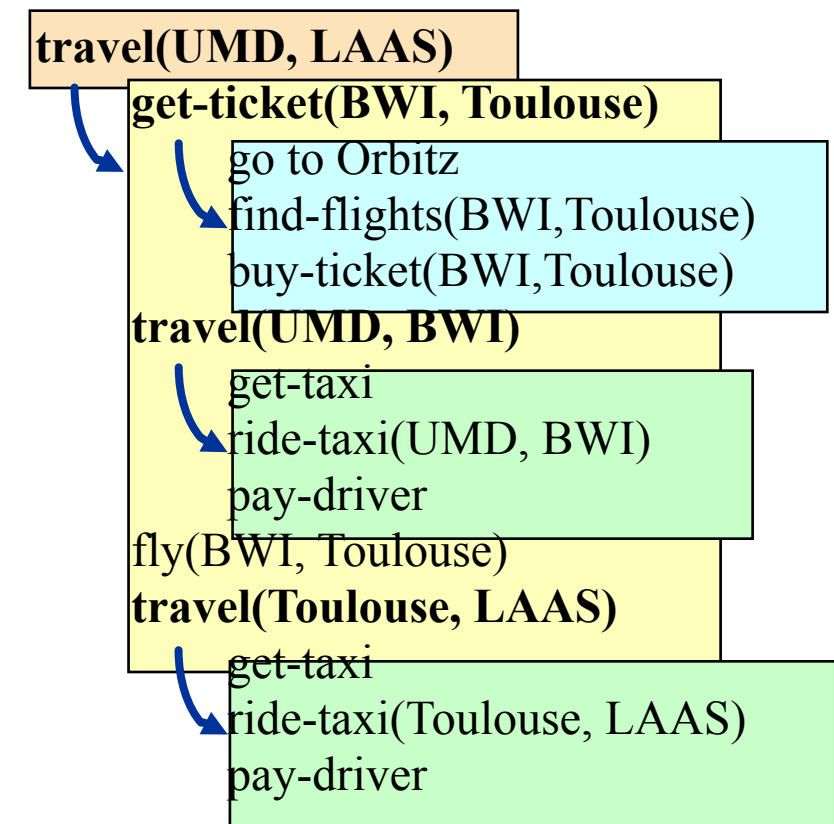
# Hierarchical Decomposition & Problem Reduction

- To get to a conference in ?x, get to the airport, take a plane to ?x, then go to the conference hotel
  - To get to the airport, either drive or take a cab
  - If you have money for the taxi fare:
  - Enter the cab, say “I want to go to ?y”, wait until you are at ?y, pay the fare, then exit the taxi”
- Idea is to capture the hierarchical structure of a planning domain
  - assuming it contains complex tasks and schemas for reducing them.
- Reduction schemas:
  - given by the designer
  - express preferred ways to accomplish a task

- Main idea behind HTN planning
- STNs: Representation and planning algorithms
  - Total order
  - Partial order
- Generalizing the formalism and algorithm to HTN
- Expressivity: comparison to classical planning and control rules
- Experimental Results

# HTN Planning

- A type of **problem reduction**
- Decompose **tasks** into **subtasks**
- Handle constraints (e.g., taxi not good for long distances)
- Resolve interactions (e.g., take taxi early enough to catch plane)
- If necessary, backtrack and try other decompositions

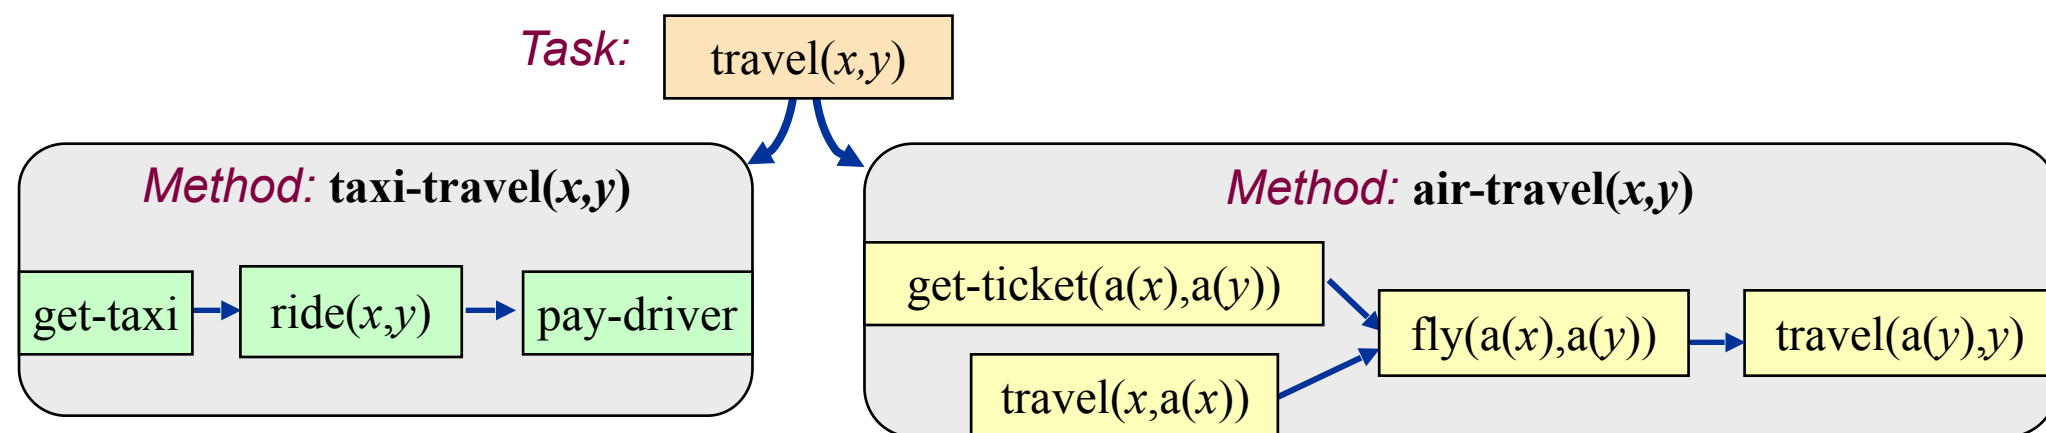


Tasks are in rectangle and methods in ovals



# HTN Planning

- HTN planners may be domain-specific
  - e.g., see Chapters 20 (robotics) and 23 (bridge)
- Or they may be domain-configurable
  - Domain-independent planning engine
  - Domain description defining operators and also methods
  - Problem description
    - domain description, initial state, initial task network



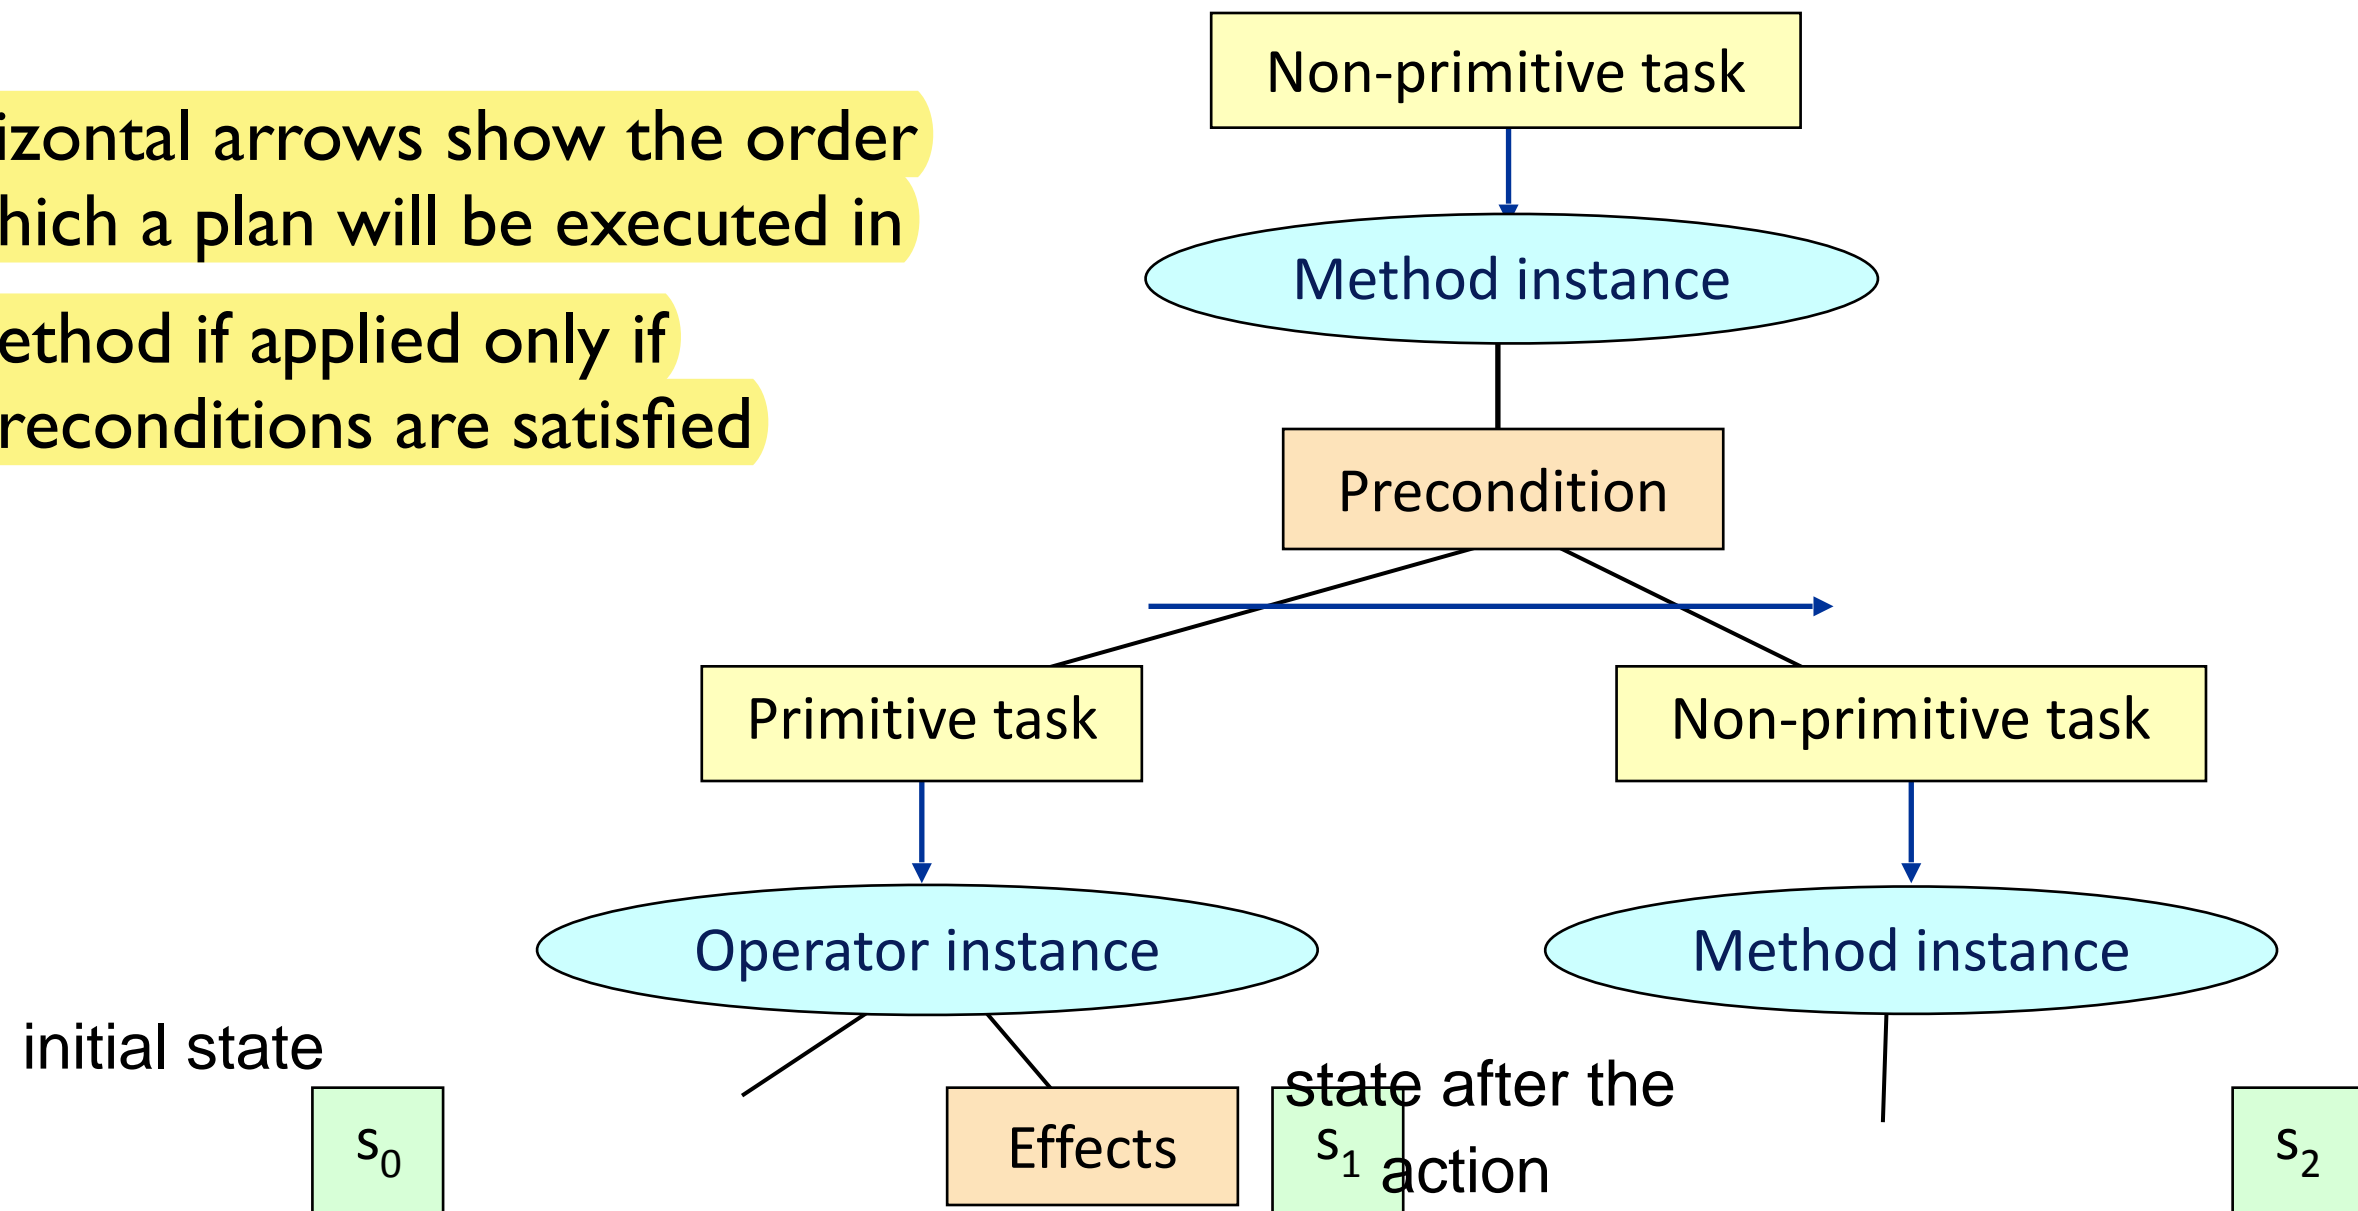
# HTN Planning

- HTN planners may be domain-specific
  - e.g., see Chapters 20 (robotics) and 23 (bridge)
- Or they may be domain-configurable
  - Domain-independent planning engine
  - Domain description
    - methods, operators
  - Problem description
    - domain description, initial state, initial task network

```
Abstract-search( $u$ )
  if Terminal( $u$ ) then return( $u$ )
   $u \leftarrow$  Refine( $u$ )           ;; refinement step
   $B \leftarrow$  Branch( $u$ )        ;; branching step
   $B' \leftarrow$  Prune( $B$ )         ;; pruning step
  if  $B' = \emptyset$  then return(failure)
  nondeterministically choose  $v \in B'$ 
  return(Abstract-search( $v$ ))
end
```

# HTN Planning: Task Networks

- Ground/unground
- Primitive/non-primitive
- Partial/total order
- Horizontal arrows show the order in which a plan will be executed in
- A method is applied only if its preconditions are satisfied



# HTN v what we've seen so far

## ■ What stays the same:

- Each state of the world is represented by a set of atoms
- Each action corresponds to a deterministic state transition
- Terms, literals, operators, actions, plans have same meaning
- E.g. (block b1) (block b2) (block b3) (block b4) (on-table b1) (on b2 b1) (clear b2) (on-table b3) (on b4 b3) (clear b4)

## ■ What's new:

- Perform a set of tasks; not achieve a set of goals
- Methods describing ways in which tasks can be performed
- Organized collections of tasks and subtasks called task networks

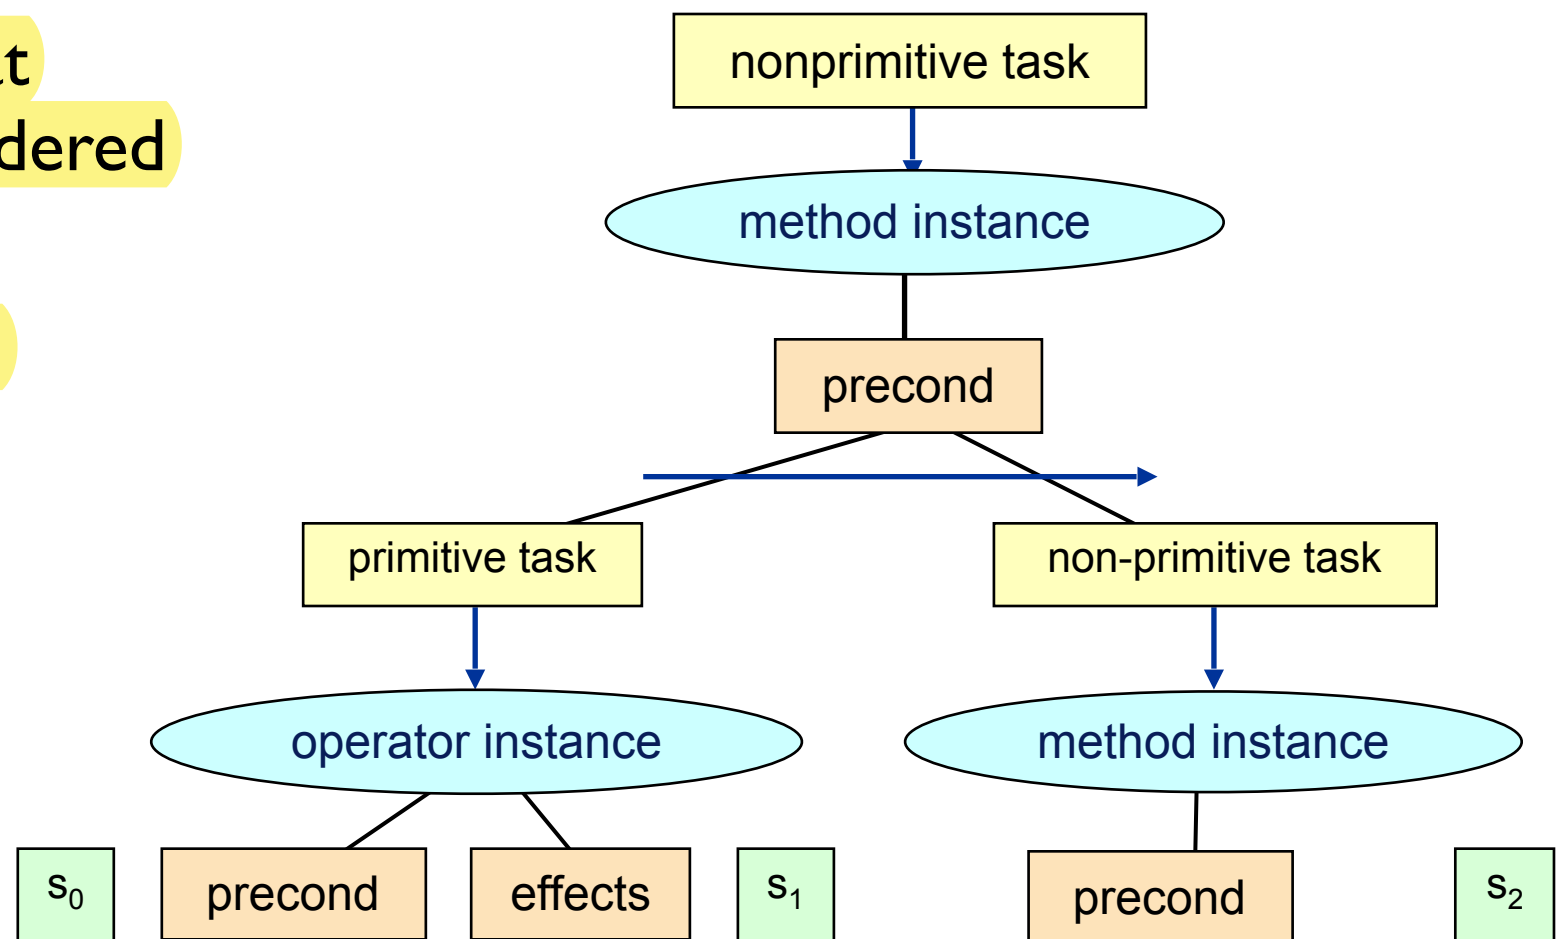
# Simple Task Network (STN) Planning

- A special case of HTN planning
- States and operators
  - The same as in classical planning
- **Task:** an expression of the form  $t(u_1, \dots, u_n)$ 
  - $t$  is a task symbol, and each  $u_i$  is a term
- Two kinds of task symbols (and tasks):
  - **Primitive:** tasks that we know how to execute directly
    - task symbol is an operator name
  - **Non-primitive:** tasks that must be decomposed into subtasks
    - use methods (next slide)

# STN: Domains, Planning Problems, Solutions

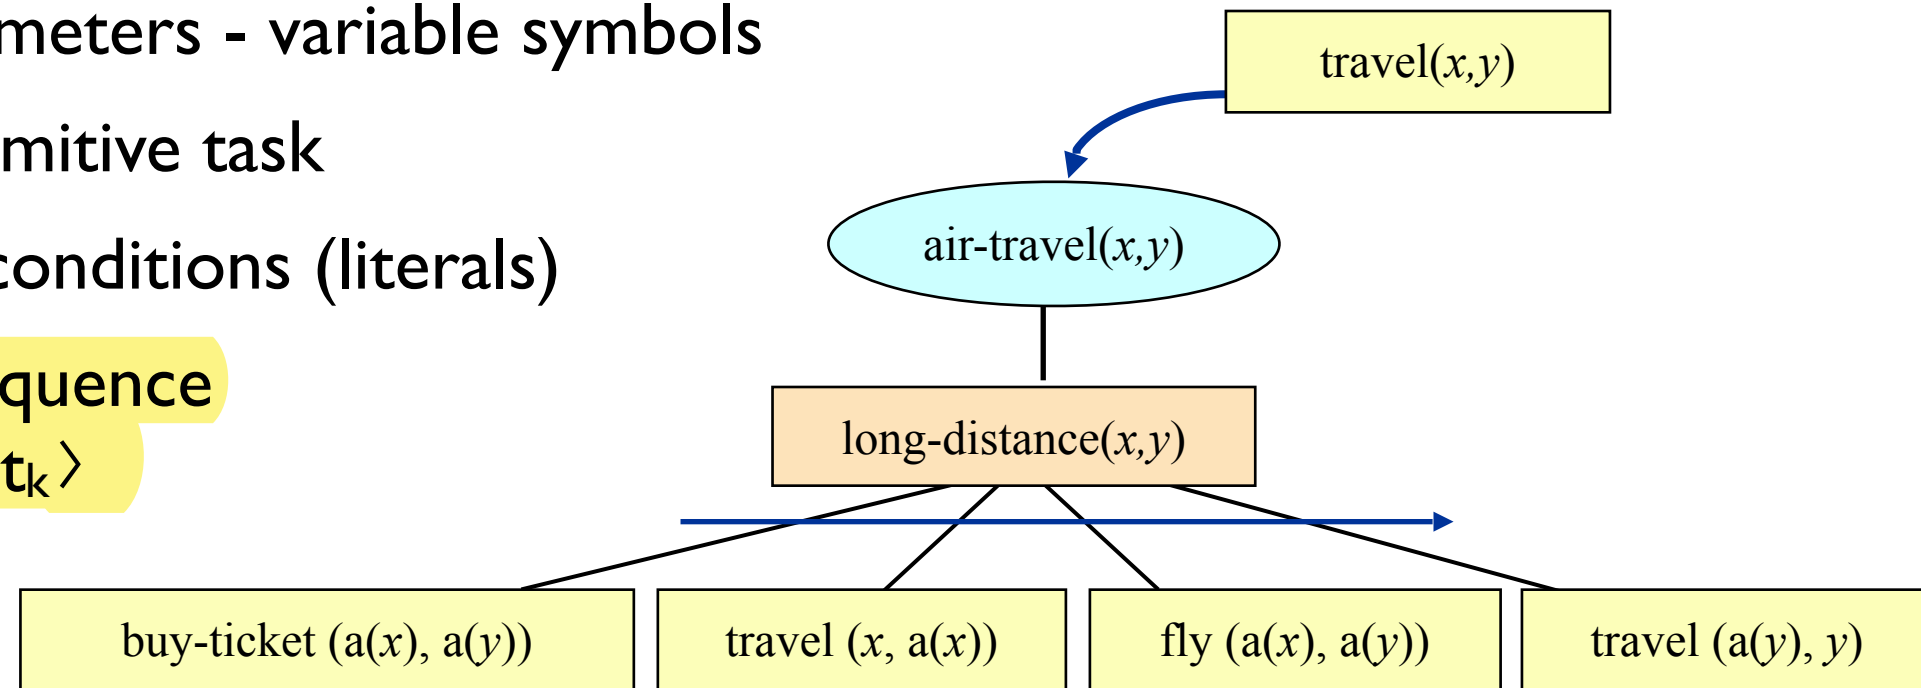
initial task network: our goal plan which consists of primitive tasks that can satisfy our goal

- Domain: methods, operators:  $D=(O,M)$
- Problem: initial state, initial task network, operators, methods:  $P=(S_0, w_j, O, M)$
- Total-order STN domain and problem:
  - Same as above except that all methods are totally ordered
- Solution: any executable plan that can be generated by recursively applying
  - methods to non-primitive tasks
  - operators to primitive tasks



# STN: Methods (totally-ordered)

- Totally-ordered method: a 4-tuple  
 $m = (\text{name}(m), \text{task}(m), \text{precond}(m), \text{subtasks}(m))$ 
  - $\text{name}(m)$ : an expression of the form  $n(x_1, \dots, x_n)$
  - $x_1, \dots, x_n$  are parameters - variable symbols
  - $\text{task}(m)$ : a nonprimitive task
  - $\text{precond}(m)$ : preconditions (literals)
  - $\text{subtasks}(m)$ : a sequence of tasks  $\langle t_1, \dots, t_k \rangle$



- $\text{air-travel}(x,y)$ 
  - task:  $\text{travel}(x,y)$
  - precondition:  $\text{long-distance}(x,y)$
  - subtasks:  $\langle \text{buy-ticket}(a(x),a(y)), \text{travel}(x,a(x)), \text{fly}(a(x),a(y)), \text{travel}(a(y),y) \rangle$

# STN: Methods (partially-ordered)

- Partially-ordered method: a 4-tuple

$m = (\text{name}(m), \text{task}(m), \text{precond}(m), \text{subtasks}(m))$

- $\text{name}(m)$ : an expression of the form  $n(x_1, \dots, x_n)$

- $x_1, \dots, x_n$  are parameters - variable symbols

- $\text{task}(m)$ : a nonprimitive task

- $\text{precond}(m)$ : preconditions (literals)

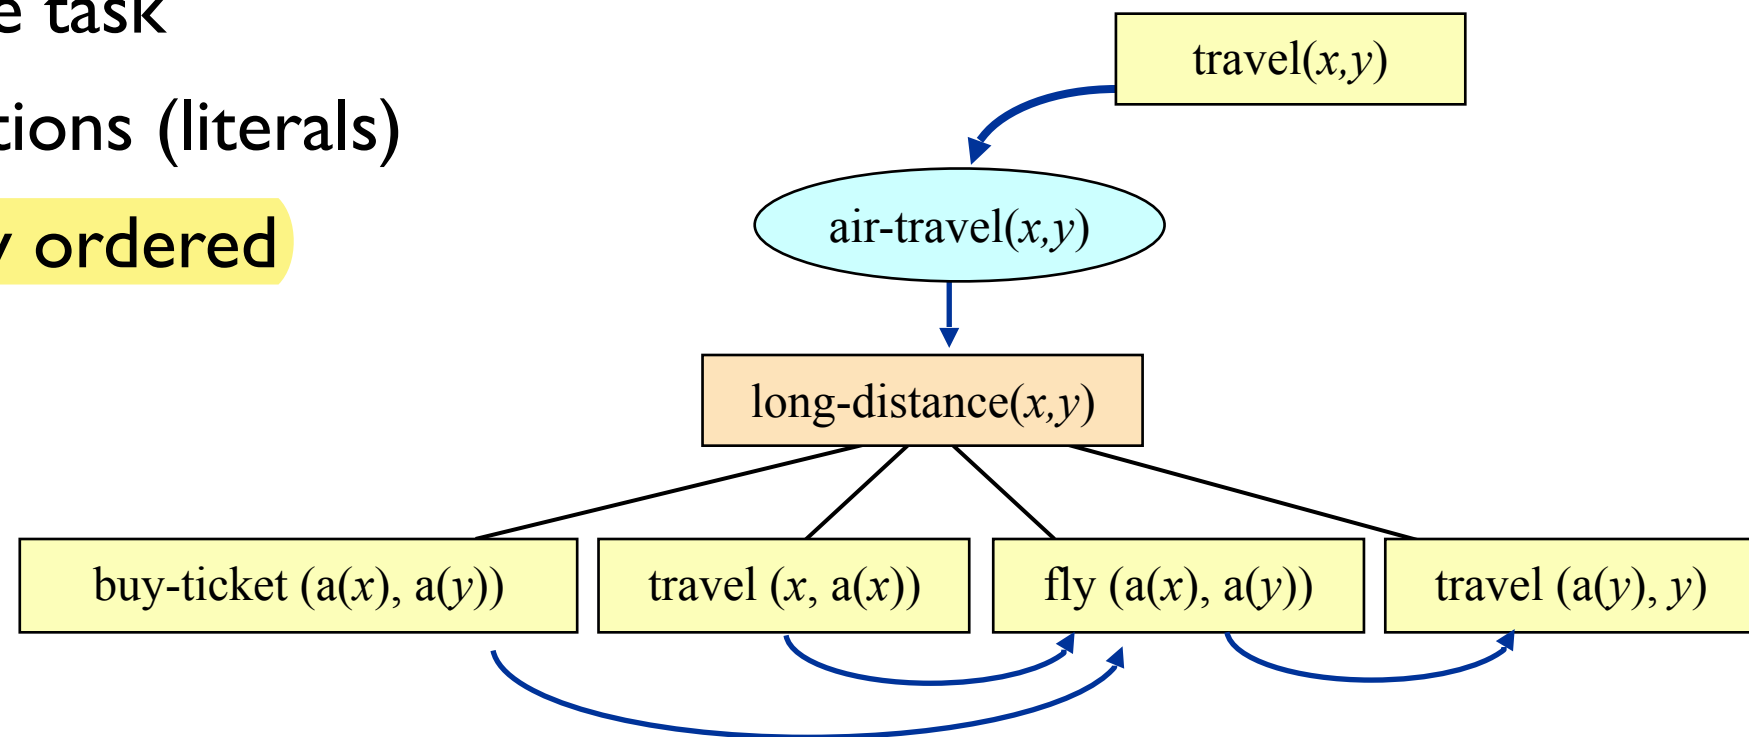
- $\text{subtasks}(m)$ : a partially ordered set of tasks  $\{t_1, \dots, t_k\}$

- $\text{air-travel}(x, y)$

- task:  $\text{travel}(x, y)$

- precondition:  $\text{long-distance}(x, y)$

- network:  $u_1 = \text{buy-ticket}(a(x), a(y))$ ,  $u_2 = \text{travel}(x, a(x))$ ,  $u_3 = \text{fly}(a(x), a(y))$ ,  $u_4 = \text{travel}(a(y), y)$ ,  $\{(u_1, u_3), (u_2, u_3), (u_3, u_4)\}$





# Example: DWR



- Task:

- Move three stacks of containers in a way that preserves the order of the containers

- One way to move each stack:


- First move the containers from p to an intermediate pile r

- Then move them from r to q



# Example: DWR

- Methods (informal):
  - move each stack twice:
    - move stack to intermediate pile (reversing order)  
and then to final destination (reversing order again)
  - move stack:
    - repeatedly/recursively move  
the topmost container  
until the stack is empty
  - move top-most:
    - take followed by put action



# Example: DWR Total-Order Formulation

take-and-put( $c, k, l_1, l_2, p_1, p_2, x_1, x_2$ ):

task: move-topmost-container( $p_1, p_2$ )

precond:  $\text{top}(c, p_1), \text{on}(c, x_1)$  ; true if  $p_1$  is not empty  
 $\text{attached}(p_1, l_1), \text{belong}(k, l_1)$  ; bind  $l_1$  and  $k$   
 $\text{attached}(p_2, l_2), \text{top}(x_2, p_2)$  ; bind  $l_2$  and  $x_2$

subtasks:  $\langle \text{take}(k, l_1, c, x_1, p_1), \text{put}(k, l_2, c, x_2, p_2) \rangle$

recursive-move( $p, q, c, x$ ):

task: move-stack( $p, q$ )

precond:  $\text{top}(c, p), \text{on}(c, x)$  ; true if  $p$  is not empty

subtasks:  $\langle \text{move-topmost-container}(p, q), \text{move-stack}(p, q) \rangle$   
*:: the second subtask recursively moves the rest of the stack*

do-nothing( $p, q$ )

task: move-stack( $p, q$ )

precond:  $\text{top}(\text{pallet}, p)$  ; true if  $p$  is empty

subtasks:  $\langle \rangle$  ; no subtasks, because we are done



move-each-twice()

task: move-all-stacks()

precond: ; no preconditions

subtasks: ; move each stack twice:

$\langle \text{move-stack}(p1a, p1b), \text{move-stack}(p1b, p1c),$   
 $\text{move-stack}(p2a, p2b), \text{move-stack}(p2b, p2c),$   
 $\text{move-stack}(p3a, p3b), \text{move-stack}(p3b, p3c) \rangle$





# STN: Solving Total-Order Planning Problems

TFD( $s, \langle t_1, \dots, t_k \rangle, O, M$ )

if  $k = 0$  then return  $\langle \rangle$  (i.e., the empty plan)

if  $t_1$  is primitive then

$active \leftarrow \{(a, \sigma) \mid a \text{ is a ground instance of an operator in } O,$   
 $\sigma \text{ is a substitution such that } a \text{ is relevant for } \sigma(t_1),$   
 $\text{and } a \text{ is applicable to } s\}$

if  $active = \emptyset$  then return failure

nondeterministically choose any  $(a, \sigma) \in active$

$\pi \leftarrow \text{TFD}(\gamma(s, a), \sigma(\langle t_2, \dots, t_k \rangle), O, M)$

if  $\pi = \text{failure}$  then return failure

else return  $a.\pi$

else if  $t_1$  is nonprimitive then

$active \leftarrow \{m \mid m \text{ is a ground instance of a method in } M,$   
 $\sigma \text{ is a substitution such that } m \text{ is relevant for } \sigma(t_1),$   
 $\text{and } m \text{ is applicable to } s\}$

if  $active = \emptyset$  then return failure

nondeterministically choose any  $(m, \sigma) \in active$

$w \leftarrow \text{subtasks}(m).\sigma(\langle t_2, \dots, t_k \rangle)$

return  $\text{TFD}(s, w, O, M)$

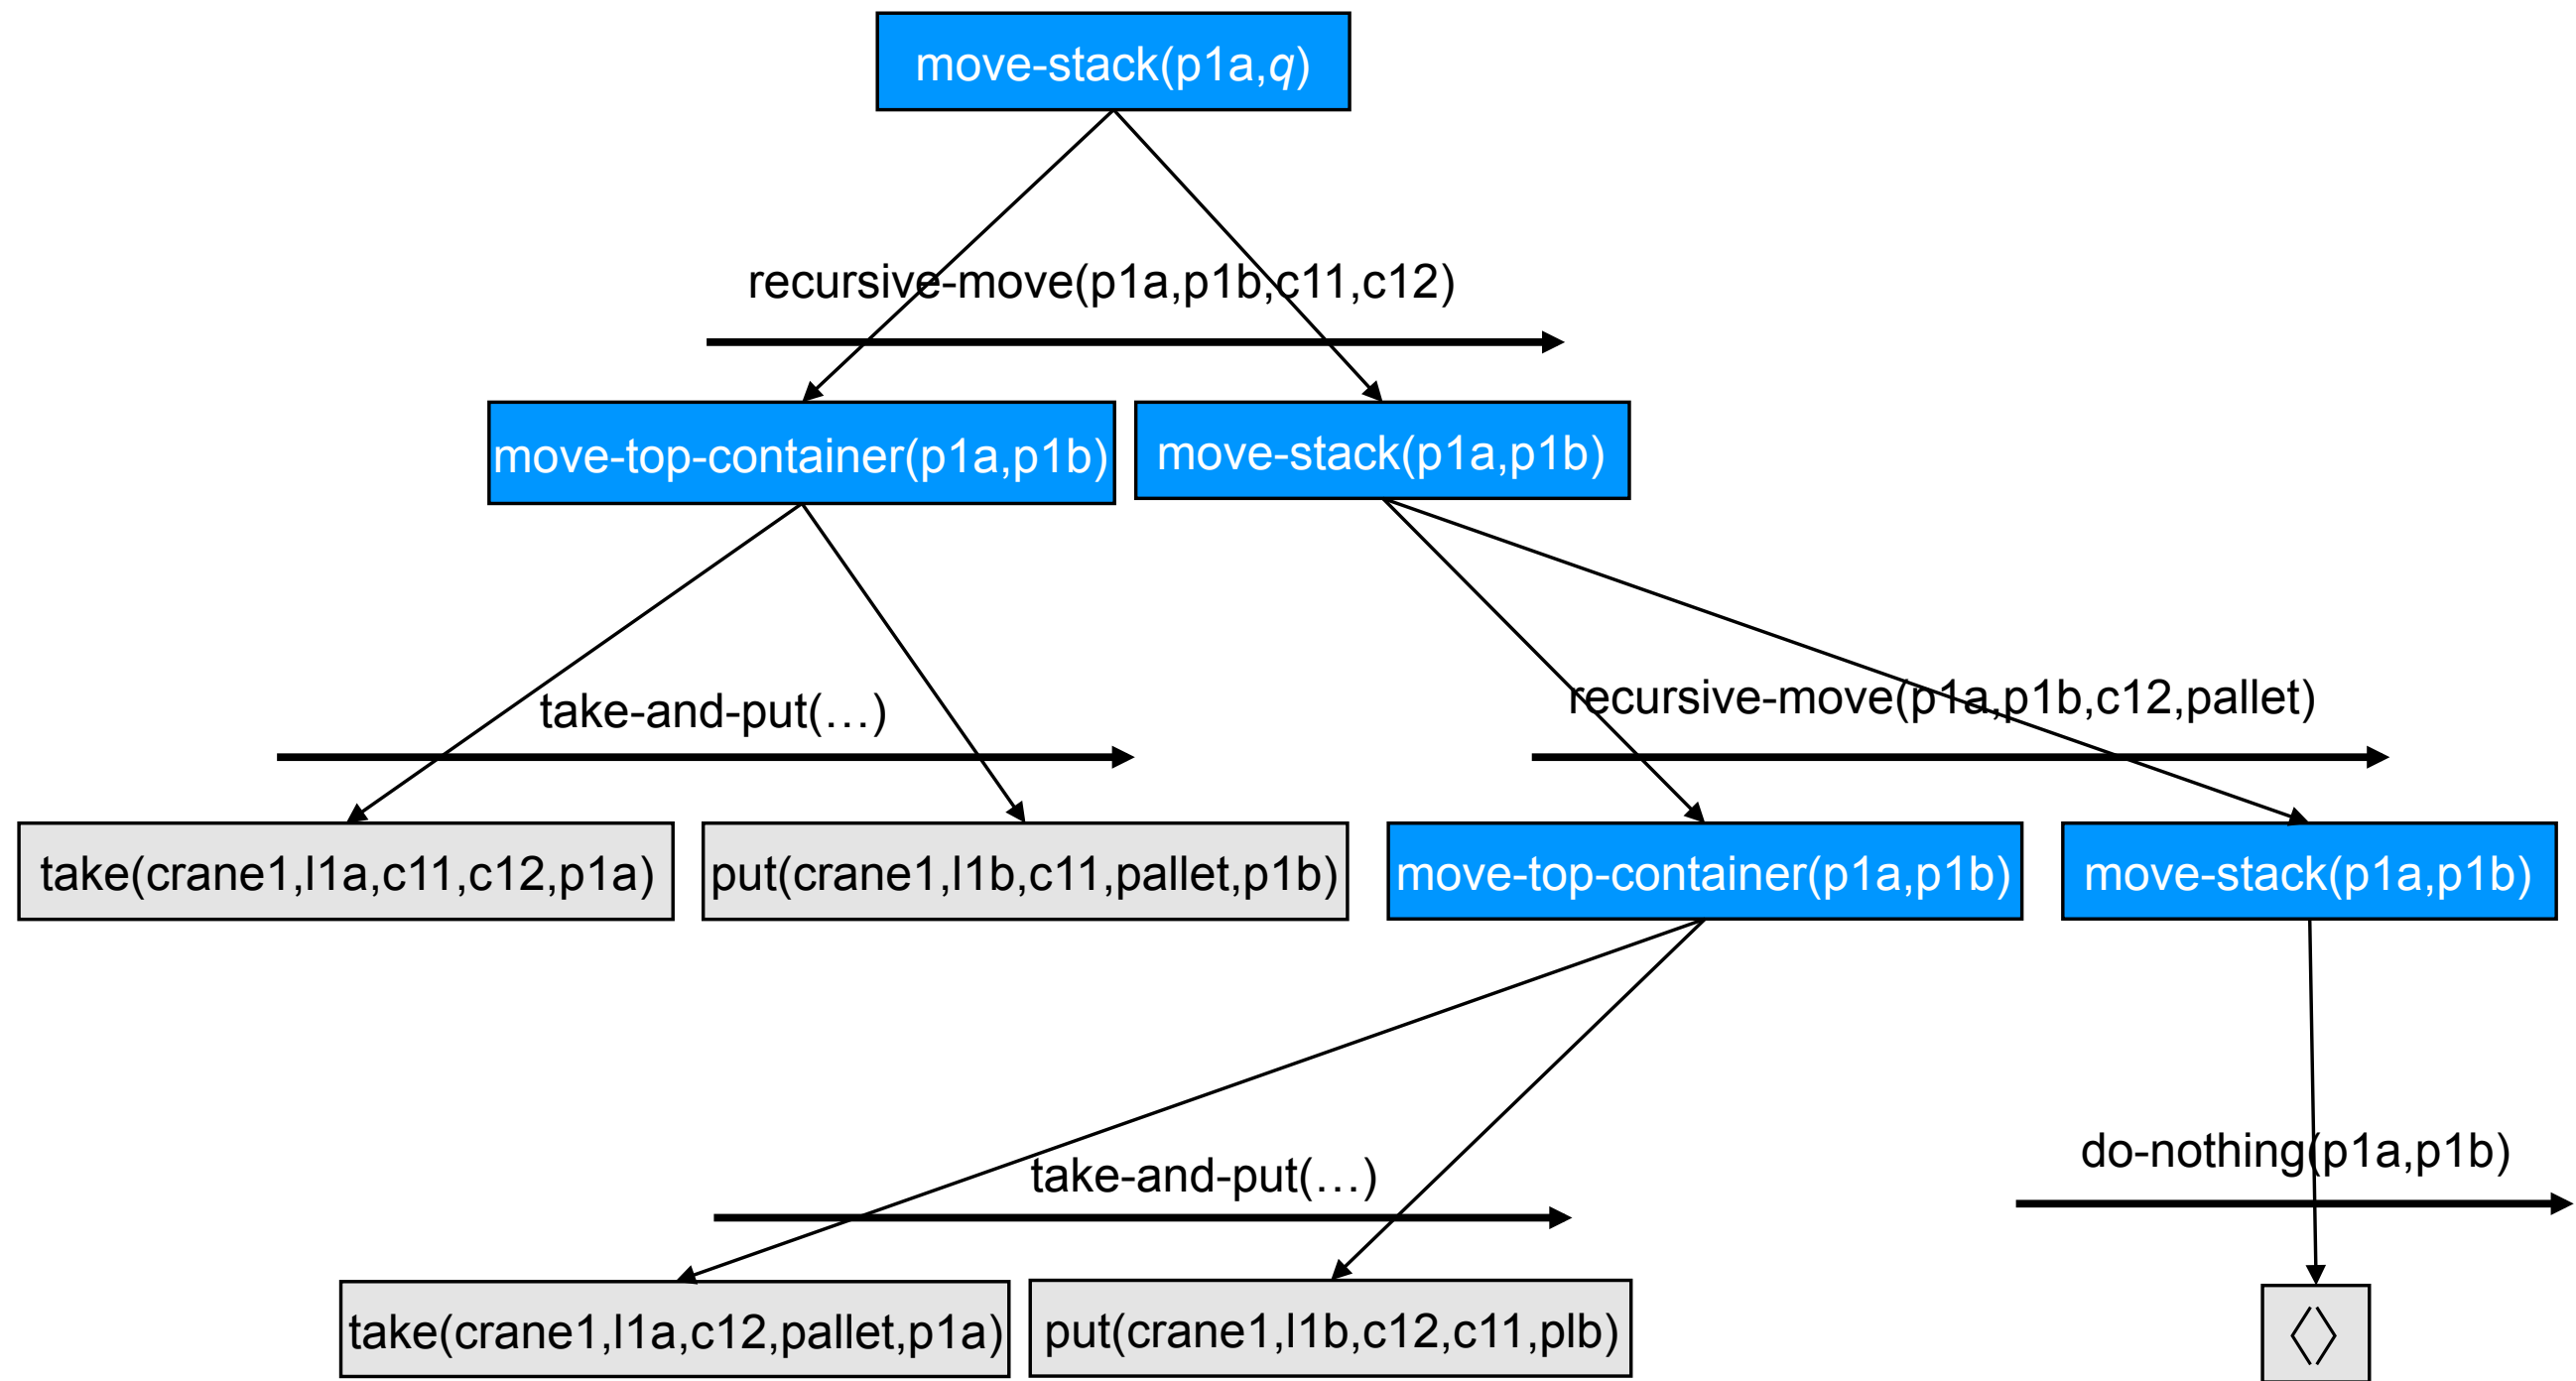
state  $s$ ; task list  $T = (\mathbf{t}_1, t_2, \dots)$   
action  $a$

state  $\gamma(s, a)$ ; task list  $T = (t_2, \dots)$

task list  $T = (\mathbf{t}_1, t_2, \dots)$   
method instance  $m$


task list  $T = (\mathbf{u}_1, \dots, \mathbf{u}_k, t_2, \dots)$

# Example: DWR Decomposition Tree - TFD



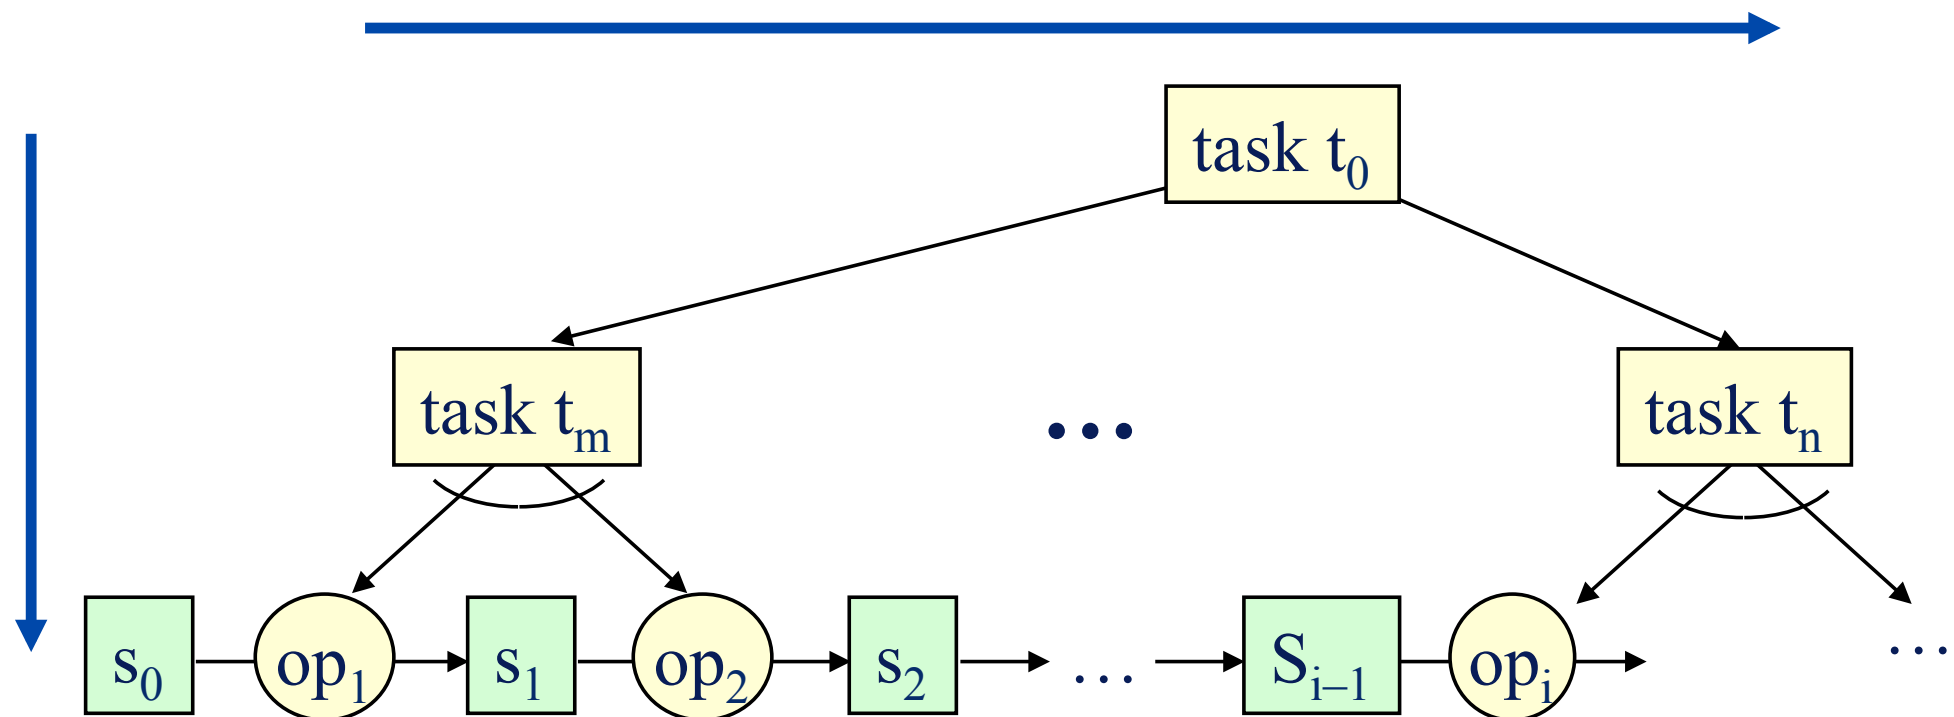
# Comparison to Forward and Backward Search

- In state-space planning, must choose whether to search forward or backward



- In HTN planning, there are two choices to make about direction:
  - forward or backward
  - up or down

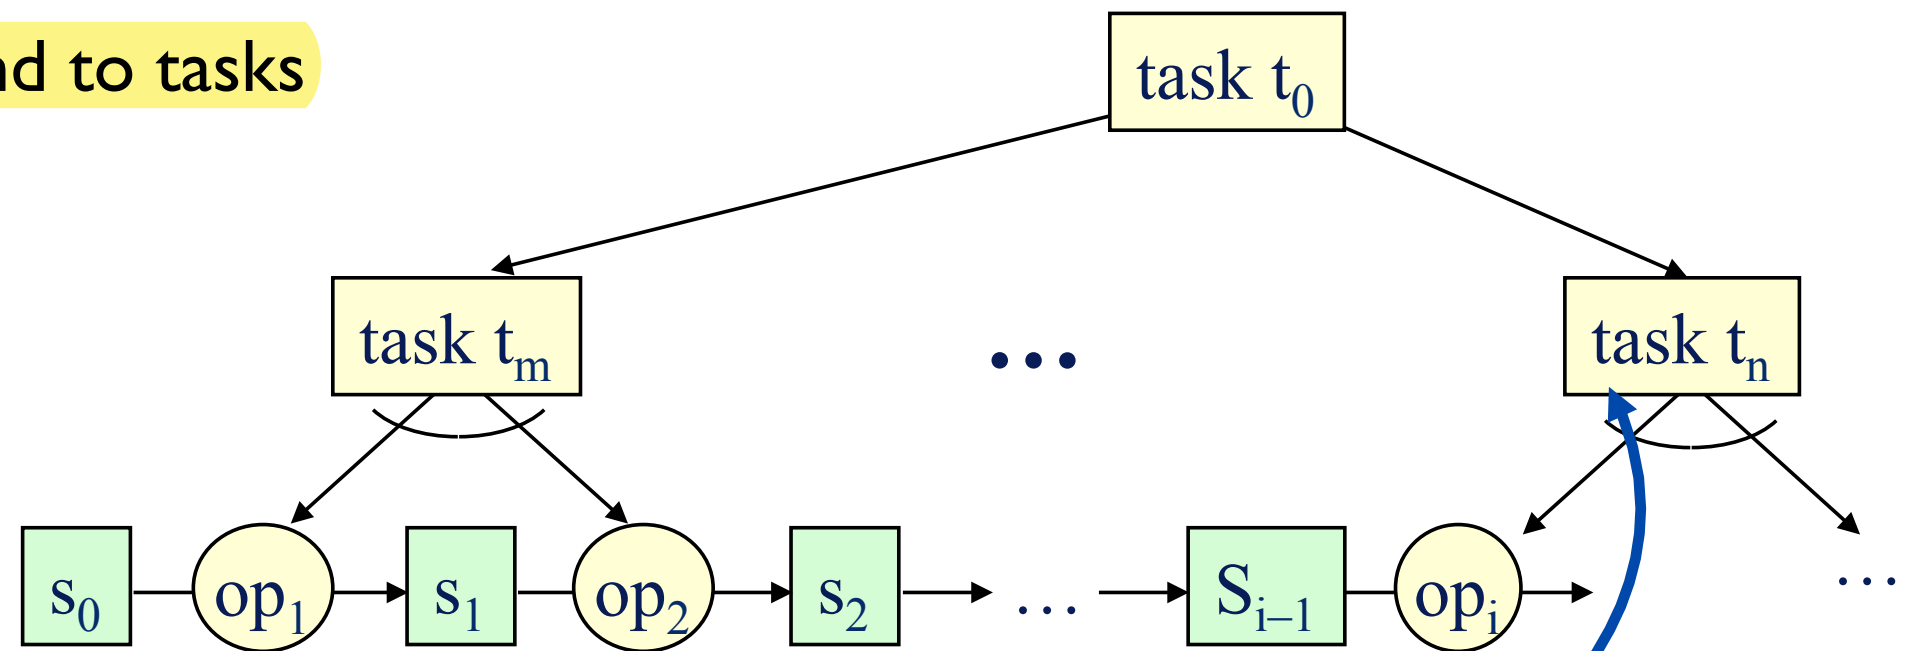
- TFD goes down and forward



# Comparison to Forward and Backward Search

- Like a backward search, TFD is goal-directed

- Goals correspond to tasks



- Like a forward search, it generates actions in the same order in which they'll be executed

- Whenever we plan the next task

- we've already planned everything that comes before it

- Thus, we know the current state of the world

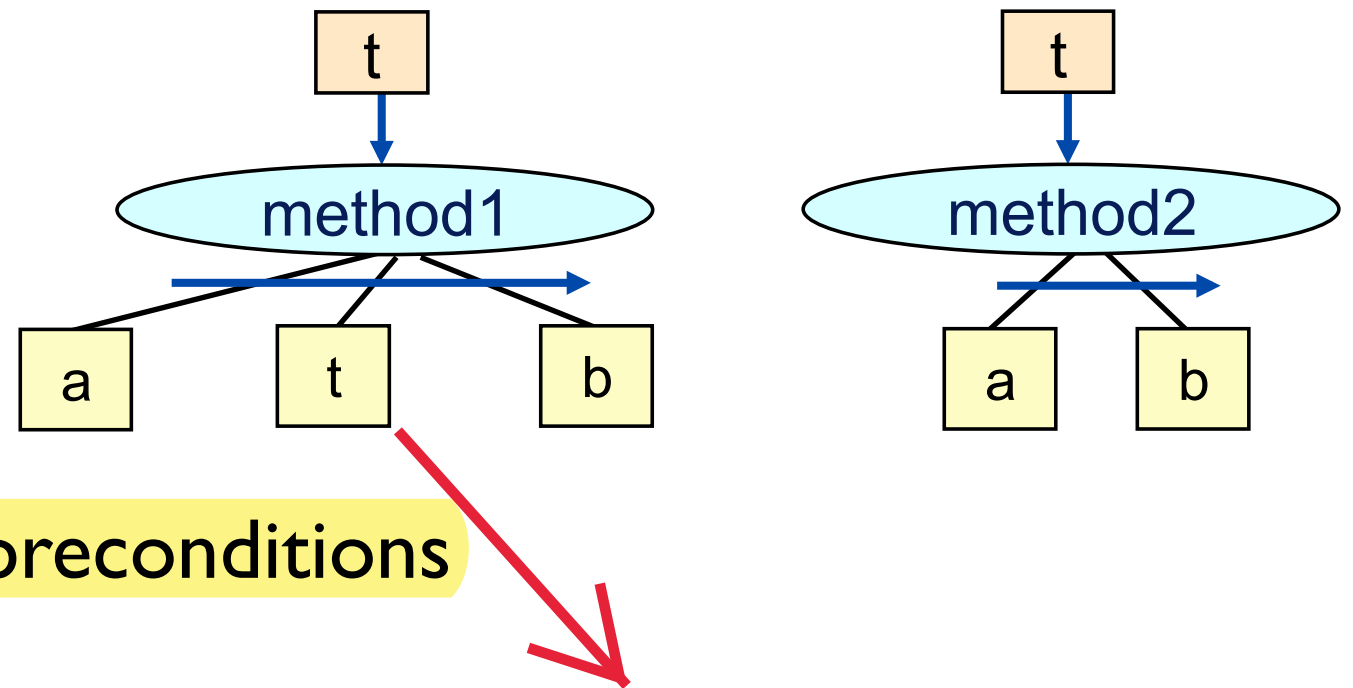
# Expressivity Relative to Classical Planning

- Any classical planning problem can be translated into an ordered-task planning problem in polynomial time
- Several ways to do this. One is roughly as follows:
  - For each goal or precondition  $e$ , create a task  $t_e$
  - For each operator  $o$  and effect  $e$ , create a method  $m_{o,e}$ 
    - Task:  $t_e$
    - Subtasks:  $t_{c1}, t_{c2}, \dots, t_{cn}, o$ , where  $c1, c2, \dots, cn$  are the preconditions of  $o$
    - Partial-ordering constraints: each  $t_{ci}$  precedes  $o$
- There are HTN planning problems that cannot be translated into classical planning problems at all
  - Example on the next page



# Example: Classical planning can not represent this

- Two methods:
  - No arguments
  - No preconditions
- Two operators, a and b
  - Again, no arguments and no preconditions



- Initial state is empty, initial task is t
- Set of solutions is  $\{a^n b^n \mid n > 0\}$

here we can have an infinite planning that is impossible in classical


- No classical planning problem has this set of solutions
  - The state transition system is a finite state automaton
  - No finite state automaton can recognise  $\{a^n b^n \mid n > 0\}$

# Increasing Expressivity Further

- Knowing the current state makes it easy to do things that would be difficult otherwise
- States can be arbitrary data structures

Us: East declarer, West dummy  
Opponents: defenders, South & North  
Contract: East – 3NT  
On lead: West at trick 3


East: ♠KJ74  
West: ♠A2  
Out: ♠QT98653



- Preconditions and effects can include
  - logical inferences (e.g., Horn clauses)
  - complex numeric computations
  - interactions with other software packages
- Example: SHOP <http://www.cs.umd.edu/projects/shop>

# Example

- Simple travel planning domain:
  - Go from one location to another
  - State-variable formulation



*method* travel-by-foot

precond:  $distance(x, y) \leq 2$

task:  $travel(a, x, y)$

subtasks:  $walk(a, x, y)$

*method* travel-by-taxi

task:  $travel(a, x, y)$

precond:  $cash(a) \geq 1.5 + 0.5 \times distance(x, y)$

subtasks:  $\langle call-taxi(a, x), ride(a, x, y), pay-driver(a, x, y) \rangle$

*operator* walk

precond:  $location(a) = x$

effects:  $location(a) \leftarrow y$

*operator* call-taxi( $a, x$ )

effects:  $location(taxi) \leftarrow x$

*operator* ride-taxi( $a, x$ )

precond:  $location(taxi) = x, location(a) = x$

effects:  $location(taxi) \leftarrow y, location(a) \leftarrow y$


*operator* pay-driver( $a, x, y$ )

precond:  $cash(a) \geq 1.5 + 0.5 \times distance(x, y)$

effects:  $cash(a) \leftarrow cash(a) - 1.5 + 0.5 \times distance(x, y)$

# Planning Problem

I am at home, I have \$20,  
I want to go to a park 8 miles away



Initial task: `travel(me,home,park)`

`travel-by-foot`

`travel-by-taxi`

Precond:  $\text{distance}(\text{home}, \text{park}) \leq 2$

Precond:  $\text{cash}(\text{me}) \geq 1.50 + 0.50 * \text{distance}(\text{home}, \text{park})$

Precondition fails

Precondition succeeds

Decomposition into subtasks

Initial state

`call-taxi(me,home)`

Precond: ...  
Effects: ...

$s_1$

`ride(me,home,park)`

Precond: ...  
Effects: ...

$s_2$

`pay-driver(me,home,park)`

Precond: ...  
Effects: ...

$s_3$

Final state

$s_0 = \{\text{location}(\text{me})=\text{home}, \text{cash}(\text{me})=20, \text{distance}(\text{home}, \text{park})=8\}$

$s_1 = \{\text{location}(\text{me})=\text{home}, \text{location}(\text{taxi})=\text{home}, \text{cash}(\text{me})=20, \text{distance}(\text{home}, \text{park})=8\}$

$s_2 = \{\text{location}(\text{me})=\text{park}, \text{location}(\text{taxi})=\text{park}, \text{cash}(\text{me})=20, \text{distance}(\text{home}, \text{park})=8\}$

$s_3 = \{\text{location}(\text{me})=\text{park}, \text{location}(\text{taxi})=\text{park}, \text{cash}(\text{me})=14.50, \text{distance}(\text{home}, \text{park})=8\}$

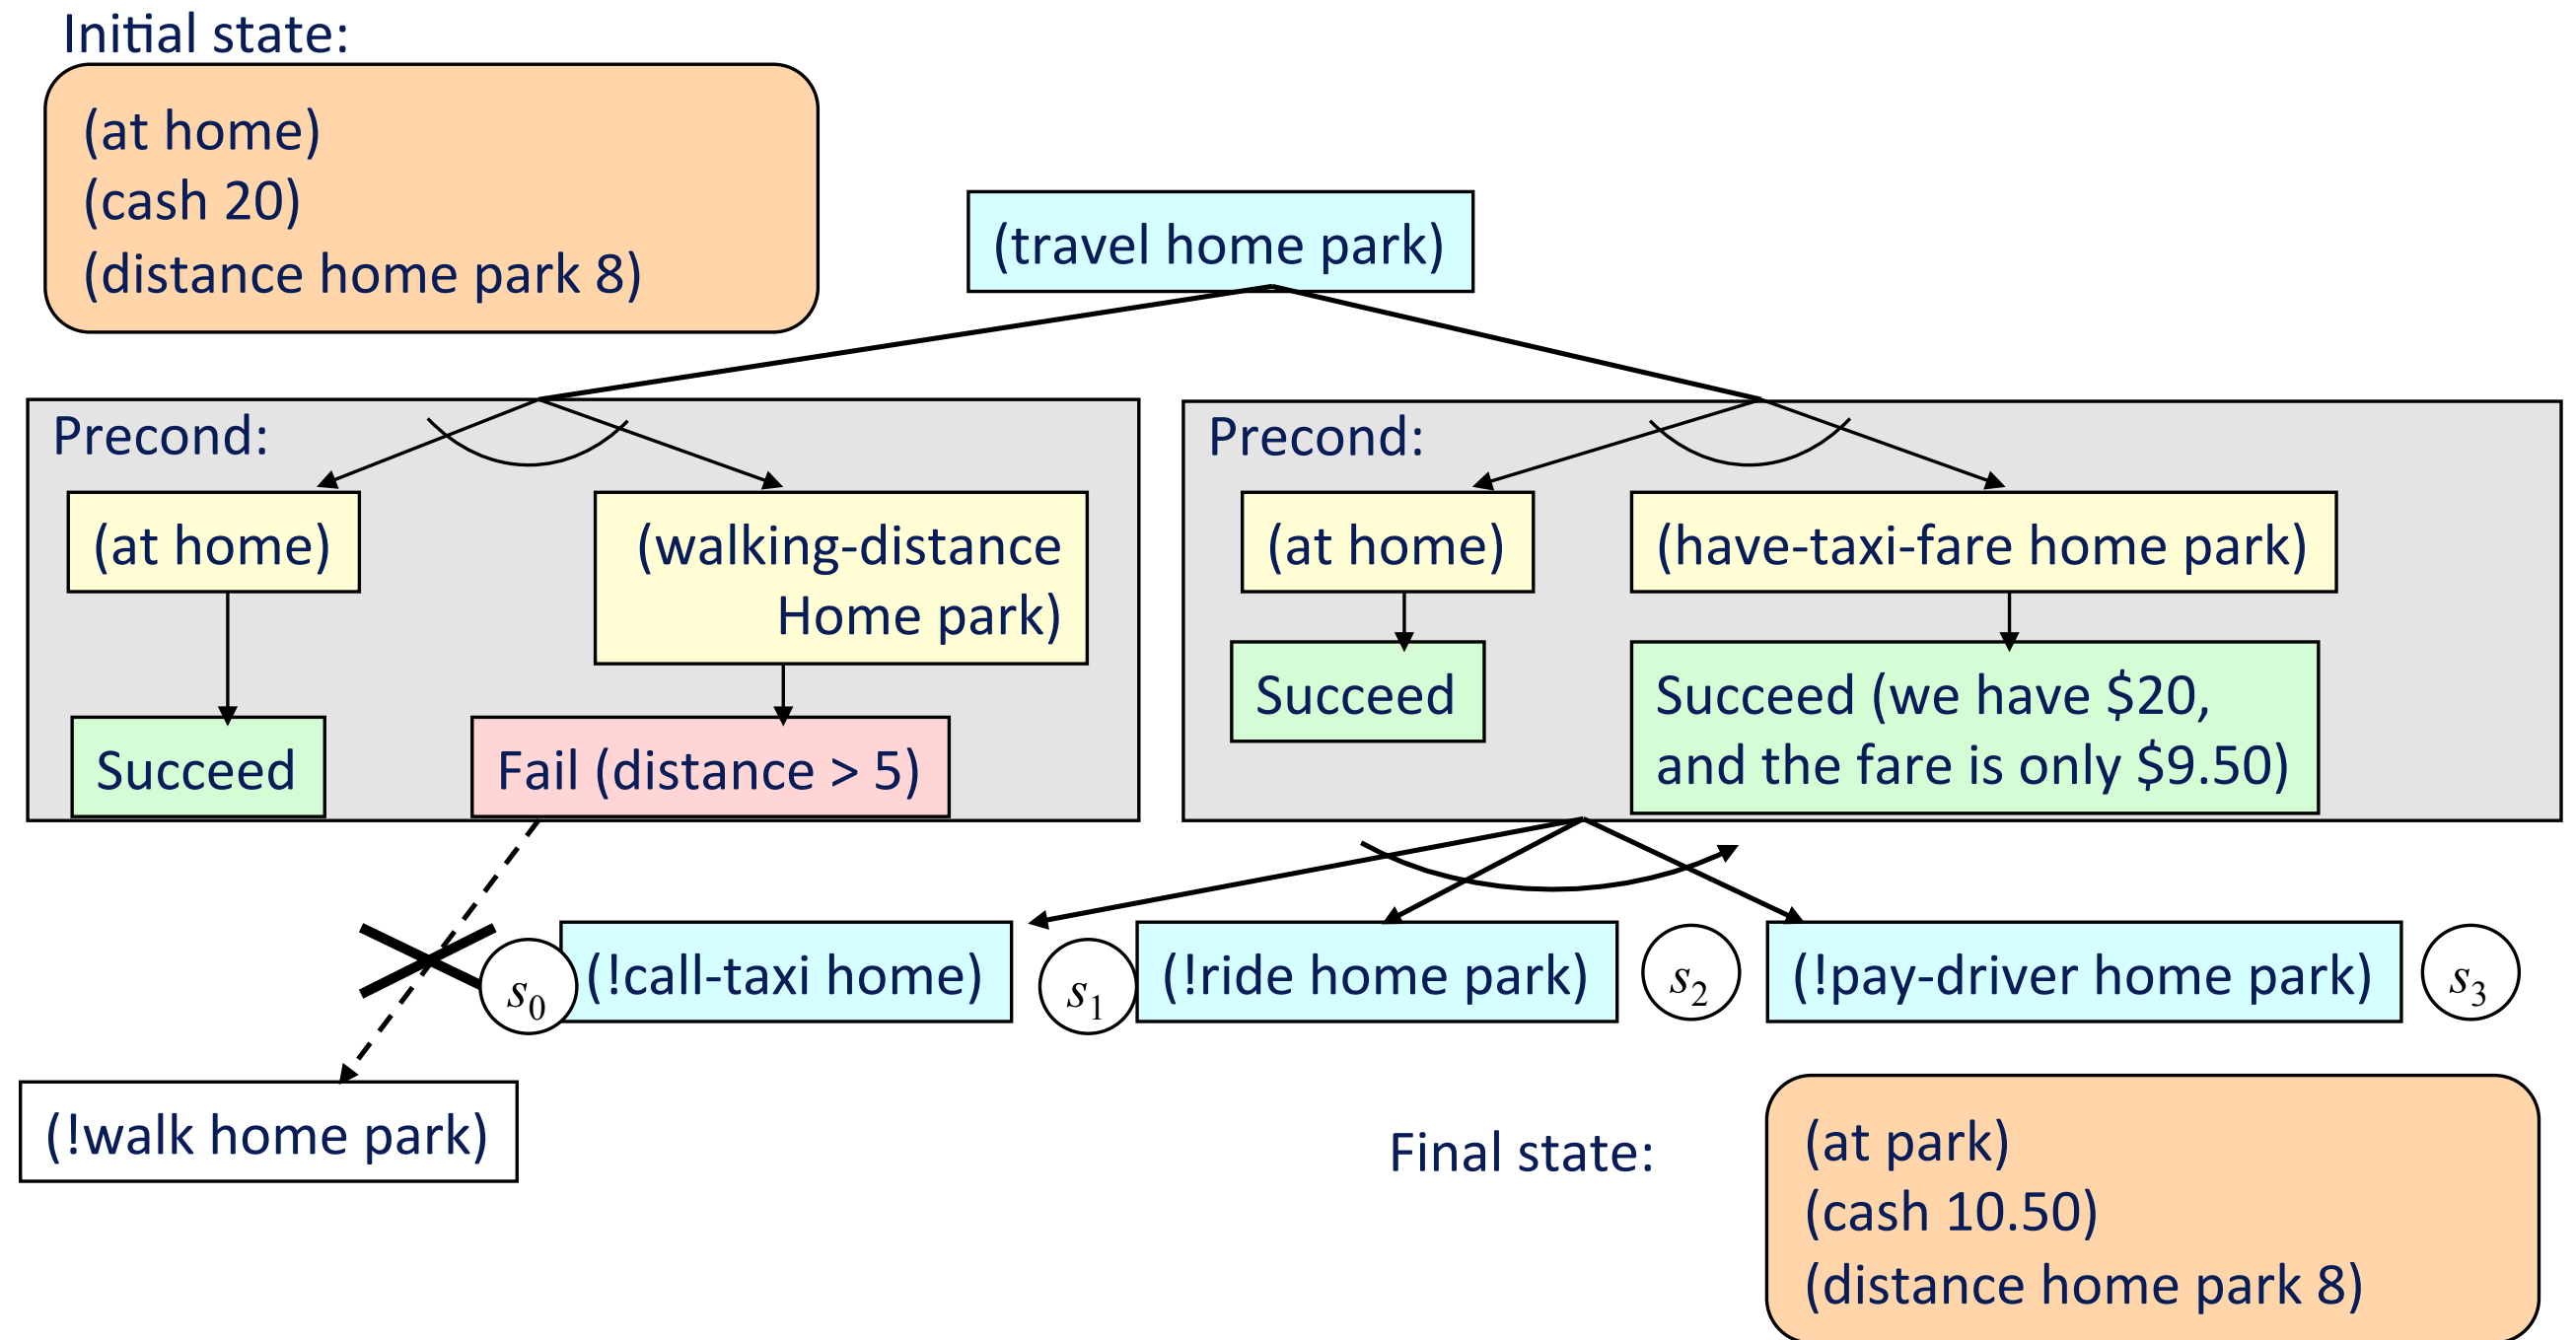
# SHOP (Simple Hierarchical Ordered Planner)

- Domain-independent algorithm for ordered task decomposition
  - Sound and complete
- Input:
  - State: a set of ground atoms
  - Task List: a linear list of tasks
  - Domain: methods, operators, axioms
- Output: one or more plans, it can return:
  - the first plan it finds
  - all possible plans
  - a least-cost plan
  - all least-cost plans

# Example: SHOP

- Initial task list: ((travel home park))
- Initial state: ((at home) (cash 20) (distance home park 8))
- Methods (task, preconditions, subtasks):
  - (:method (travel ?x ?y)  
((at x) (walking-distance ?x ?y)) ' (!walk ?x ?y)) I)
  - (:method (travel ?x ?y)  
((at ?x) (have-taxi-fare ?x ?y))  
' (!call-taxi ?x) (!ride ?x ?y) (!pay-driver ?x ?y)) I)
- Axioms:
  - (:- (walking-dist ?x ?y) ((distance ?x ?y ?d) (eval (<= ?d 5))))
  - (:- (have-taxi-fare ?x ?y)  
((have-cash ?c) (distance ?x ?y ?d) (eval (>= ?c (+ 1.50 ?d)))))
- Primitive operators (task, delete list, add list)
  - (:operator (!walk ?x ?y) ((at ?x)) ((at ?y)))
  - ...

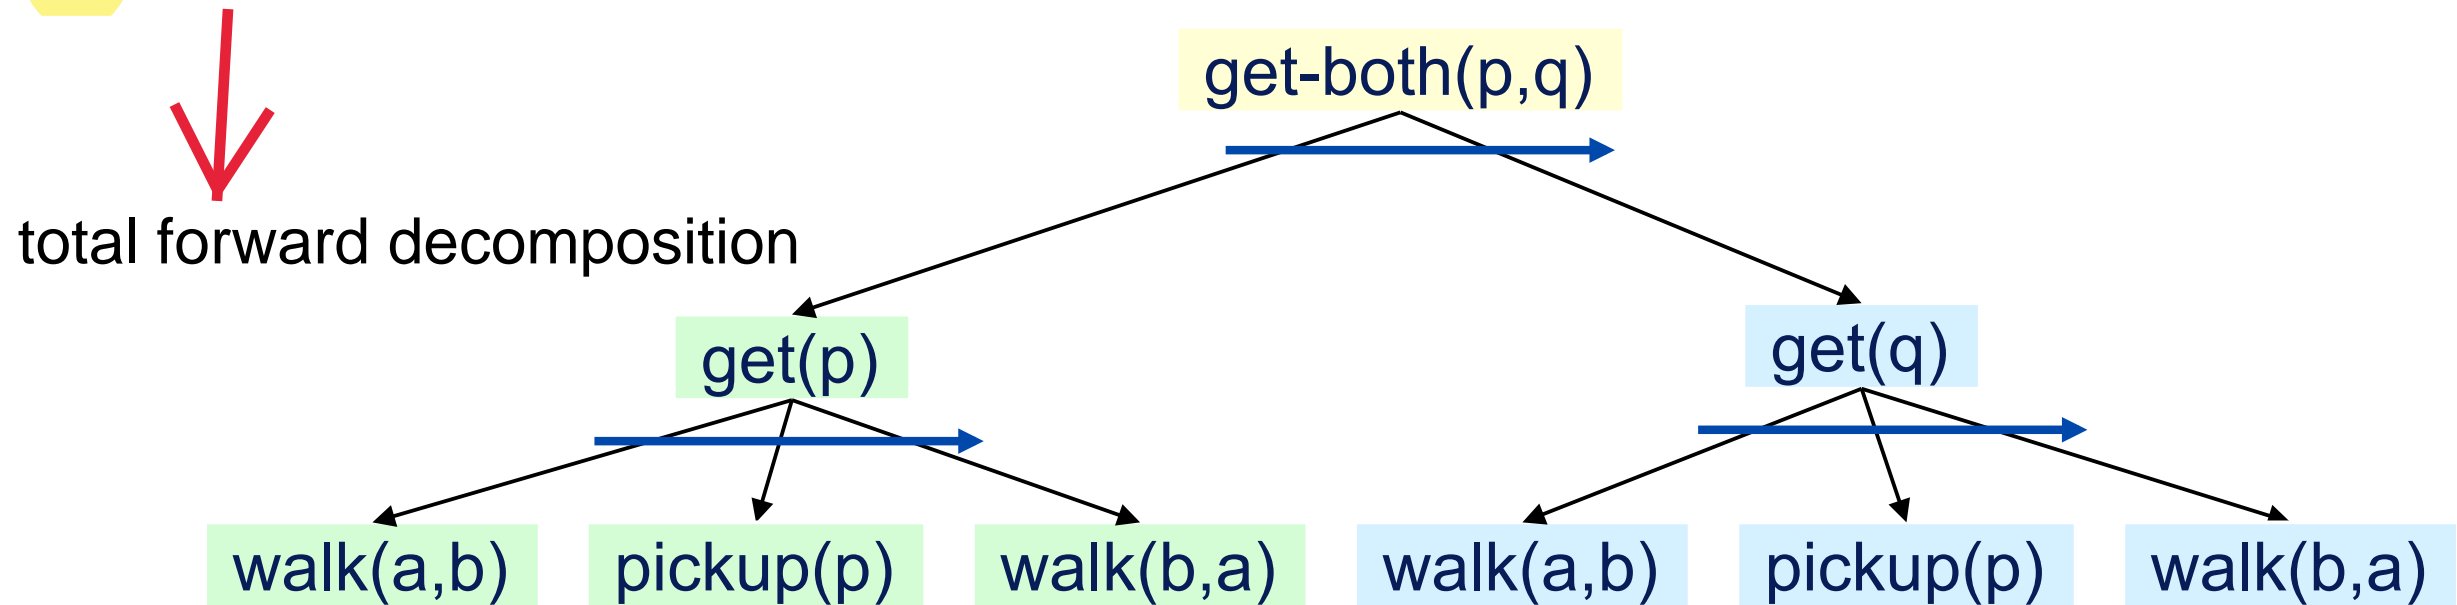
# Example: SHOP (Continued)





# Limitation of Ordered-Task Planning

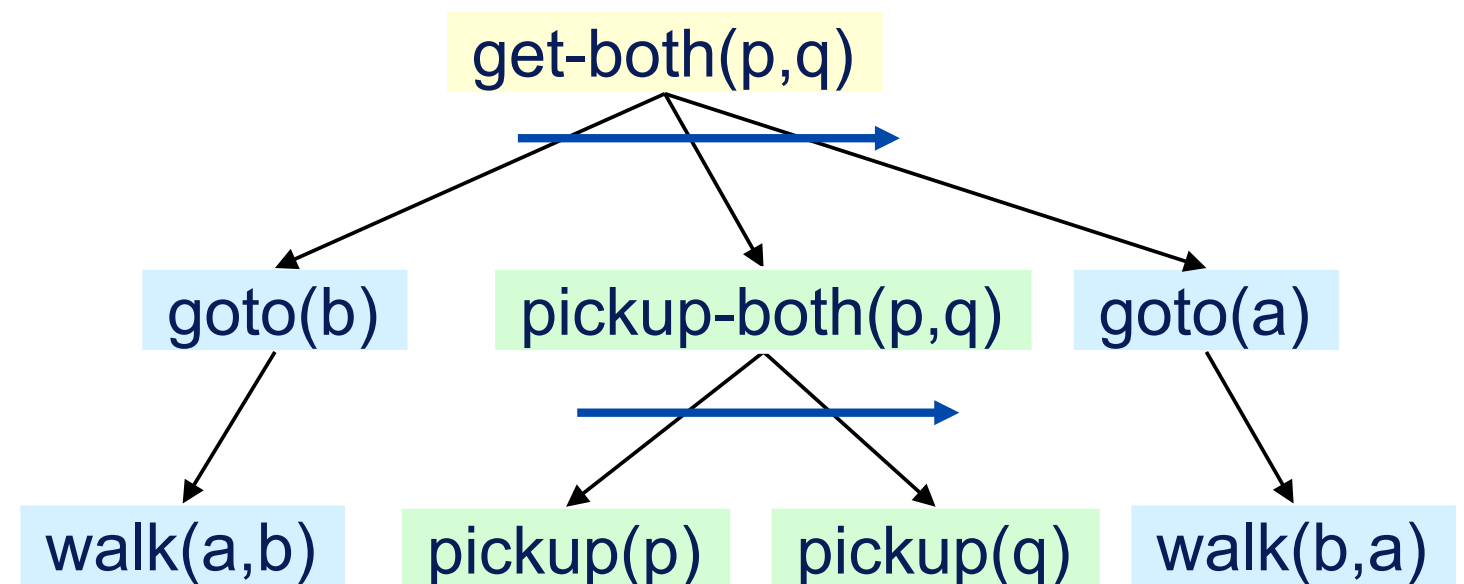
- TFD requires totally ordered methods



- Can't interleave subtasks of different tasks

- Sometimes this makes things awkward:

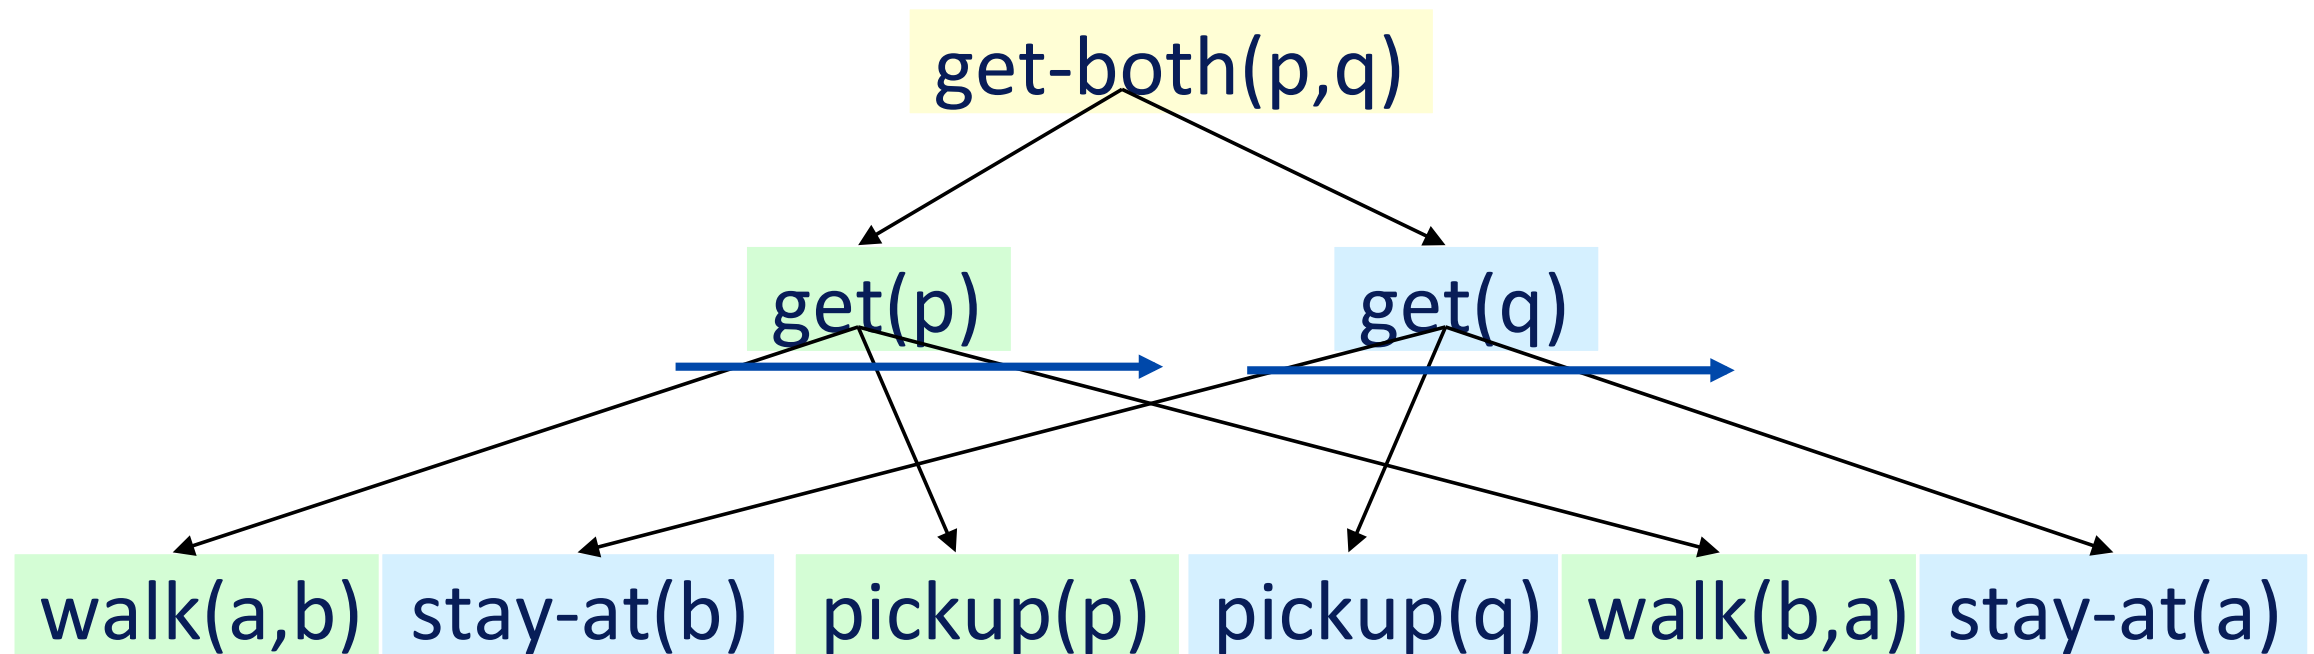
- Need to write methods that reason globally instead of locally





# Generalize TFD to interleave subtasks

- Generalize methods to allow the subtasks to be partially ordered
- Consequence: plans may interleave subtasks of different tasks



- This makes the planning algorithm more complicated

# Example: DWR Partial-Order Formulation

take-and-put( $c, k, l_1, l_2, p_1, p_2, x_1, x_2$ ):

task: move-topmost-container( $p_1, p_2$ )

precond:  $\text{top}(c, p_1), \text{on}(c, x_1)$  ; true if  $p_1$  is not empty  
 $\text{attached}(p_1, l_1), \text{belong}(k, l_1)$  ; bind  $l_1$  and  $k$   
 $\text{attached}(p_2, l_2), \text{top}(x_2, p_2)$  ; bind  $l_2$  and  $x_2$

subtasks:  $\langle \text{take}(k, l_1, c, x_1, p_1), \text{put}(k, l_2, c, x_2, p_2) \rangle$

recursive-move( $p, q, c, x$ ):

task: move-stack( $p, q$ )

precond:  $\text{top}(c, p), \text{on}(c, x)$  ; true if  $p$  is not empty

subtasks:  $\langle \text{move-topmost-container}(p, q), \text{move-stack}(p, q) \rangle$   
 ;; the second subtask recursively moves the rest of the stack

do-nothing( $p, q$ )

task: move-stack( $p, q$ )

precond:  $\text{top}(\text{pallet}, p)$  ; true if  $p$  is empty

subtasks:  $\langle \rangle$  ; no subtasks, because we are done



move-each-twice()

task: move-all-stacks()

precond: ; no preconditions

network: ; move each stack twice:

$u_1 = \text{move-stack}(p1a, p1b), u_2 = \text{move-stack}(p1b, p1c),$   
 $u_3 = \text{move-stack}(p2a, p2b), u_4 = \text{move-stack}(p2b, p2c),$   
 $u_5 = \text{move-stack}(p3a, p3b), u_6 = \text{move-stack}(p3b, p3c),$   
 $\{(u_1, u_2), (u_3, u_4), (u_5, u_6)\}$



# Solving Partial-Order STNs

$\text{PFD}(s, w, O, M)$

if  $w = \emptyset$  then return the empty plan

nondeterministically choose any  $u \in w$  that has no predecessors in  $w$

if  $t_u$  is a primitive task then

$active \leftarrow \{(a, \sigma) \mid a \text{ is a ground instance of an operator in } O,$   
 $\sigma \text{ is a substitution such that } \text{name}(a) = \sigma(t_u),$   
 $\text{and } a \text{ is applicable to } s\}$

if  $active = \emptyset$  then return failure

nondeterministically choose any  $(a, \sigma) \in active$

$\pi \leftarrow \text{PFD}(\gamma(s, a), \sigma(w - \{u\}), O, M)$

if  $\pi = \text{failure}$  then return failure

else return  $a. \pi$

else

$active \leftarrow \{(m, \sigma) \mid m \text{ is a ground instance of a method in } M,$   
 $\sigma \text{ is a substitution such that } \text{name}(m) = \sigma(t_u),$   
 $\text{and } m \text{ is applicable to } s\}$

if  $active = \emptyset$  then return failure

nondeterministically choose any  $(m, \sigma) \in active$

nondeterministically choose any task network  $w' \in \delta(w, u, m, \sigma)$

return( $\text{PFD}(s, w', O, M)$ )

$\pi = \{a_1, \dots, a_k\}; w = \{t_1, t_2, t_3, \dots\}$   
 operator instance  $a$

$\pi = \{a_1, \dots, a_k, a\}; w' = \{t_2, t_3, \dots\}$

$w = \{t_1, t_2, \dots\}$   
 method instance  $m$

$w' = \{u_1, \dots, u_k, t_2, \dots\}$



# Solving Partial-Order STNs

PFD( $s, w, O, M$ )

if  $w = \emptyset$  then return the empty plan

- Intuitively,  $w$  is a partially ordered set of tasks  $\{t_1, t_2, \dots\}$ 
  - But  $w$  may contain a task more than once
    - e.g., travel from UMD to LAAS twice
  - The mathematical definition of a set doesn't allow this
- Define  $w$  as a partially ordered set of *task nodes*  $\{u_1, u_2, \dots\}$ 
  - Each task node  $u$  corresponds to a task  $t_u$
- In my explanations, I talk about  $t$  and ignore  $u$

else return  $a.\pi$

else

$active \leftarrow \{(m, \sigma) \mid m \text{ is a ground instance of a method in } M, \\ \sigma \text{ is a substitution such that } name(m) = \sigma(t_u), \\ \text{and } m \text{ is applicable to } s\}$

if  $active = \emptyset$  then return failure

nondeterministically choose any  $(m, \sigma) \in active$

nondeterministically choose any task network  $w' \in \delta(w, u, m, \sigma)$

return(PFD( $s, w', O, M$ ))

tasks in  $w$

$O,$   
 $\sigma(t_u),$

$\pi = \{a_1, \dots, a_k\}; w = \{t_1, t_2, t_3, \dots\}$   
operator instance  $a$

$\pi = \{a_1, \dots, a_k, a\}; w' = \{t_2, t_3, \dots\}$

$w = \{t_1, t_2, \dots\}$   
method instance  $m$

$w' = \{u_1, \dots, u_k, t_2, \dots\}$

# Solving Partial-Order STNs

PFD( $s, w, O, M$ )

if  $w = \emptyset$  then return the empty plan

nondeterministically choose any  $u \in w$  that has no predecessors in  $w$

if  $t_u$  is a primitive task then

$active \leftarrow \{(a, \sigma) \mid a \text{ is a ground instance of an operator in } O,$   
 $\sigma \text{ is a substitution such that } name(a) = \sigma(t_u),$

$\delta(w, u, m, \sigma)$  has a complicated definition in the book. Here's what it means:

- We non-deterministically selected  $t_1$  as the task to do first
- Must do  $t_1$ 's first subtask before the first subtask of every  $t_i \neq t_1$
- Insert ordering constraints to ensure that this happens

$\pi = \{a_1, \dots, a_k\}; w = \{ \mathbf{t}_1, t_2, t_3 \dots \}$   
 operator instance  $\mathbf{a}$

$\pi = \{a_1 \dots, a_k, \mathbf{a}\}; w' = \{t_2, t_3 \dots\}$

$w = \{ \mathbf{t}_1, t_2, \dots \}$   
 method instance  $\mathbf{m}$

$w' = \{ \mathbf{u}_1, \dots, \mathbf{u}_k, t_2, \dots \}$

$active \leftarrow \{(m, \sigma) \mid m \text{ is a ground instance of a method in } M,$   
 $\sigma \text{ is a substitution such that } name(m) = \sigma(t_u),$   
 and  $m \text{ is applicable to } s\}$

if  $active = \emptyset$  then return failure

nondeterministically choose any  $(m, \sigma) \in active$

nondeterministically choose any task network  $w' \in \delta(w, u, m, \sigma)$

return(PFD( $s, w', O, M$ ))

# STN Summary

- PFD is sound and complete
- STN – simplified version of HTN
  - TFD – Total-order Forward Decomposition (used in SHOP)
    - Input: tasks are totally ordered
    - Output: totally ordered plan
  - PFD – Partial-order Forward Decomposition (SHOP2)
    - Input: tasks are partially ordered
    - Output: totally ordered plan
- SHOP2:
  - Won one of the top four awards in the AIPS-2002 Planning Competition
  - Freeware, open source
  - Implementation available at <http://www.cs.umd.edu/projects/shop>

- HTN – generalization of STN
  - More freedom about how to construct the task networks.
  - Can use other decomposition procedures not just forward-decomposition.
  - Like Partial-Order Planning combined with STN
    - Input: Partial-order tasks
    - Output: The resulting plan is partially ordered
  - Plans can be totally ordered or partially ordered
  - Can have constraints associated with tasks and methods
  - Things that must be true before a state, in between two given states, or after a state (replaces STN preconditions)
  - Some algorithms use causal links and threats like those in PSP

# TLPlan's Expressivity Compared with SHOP and SHOP2

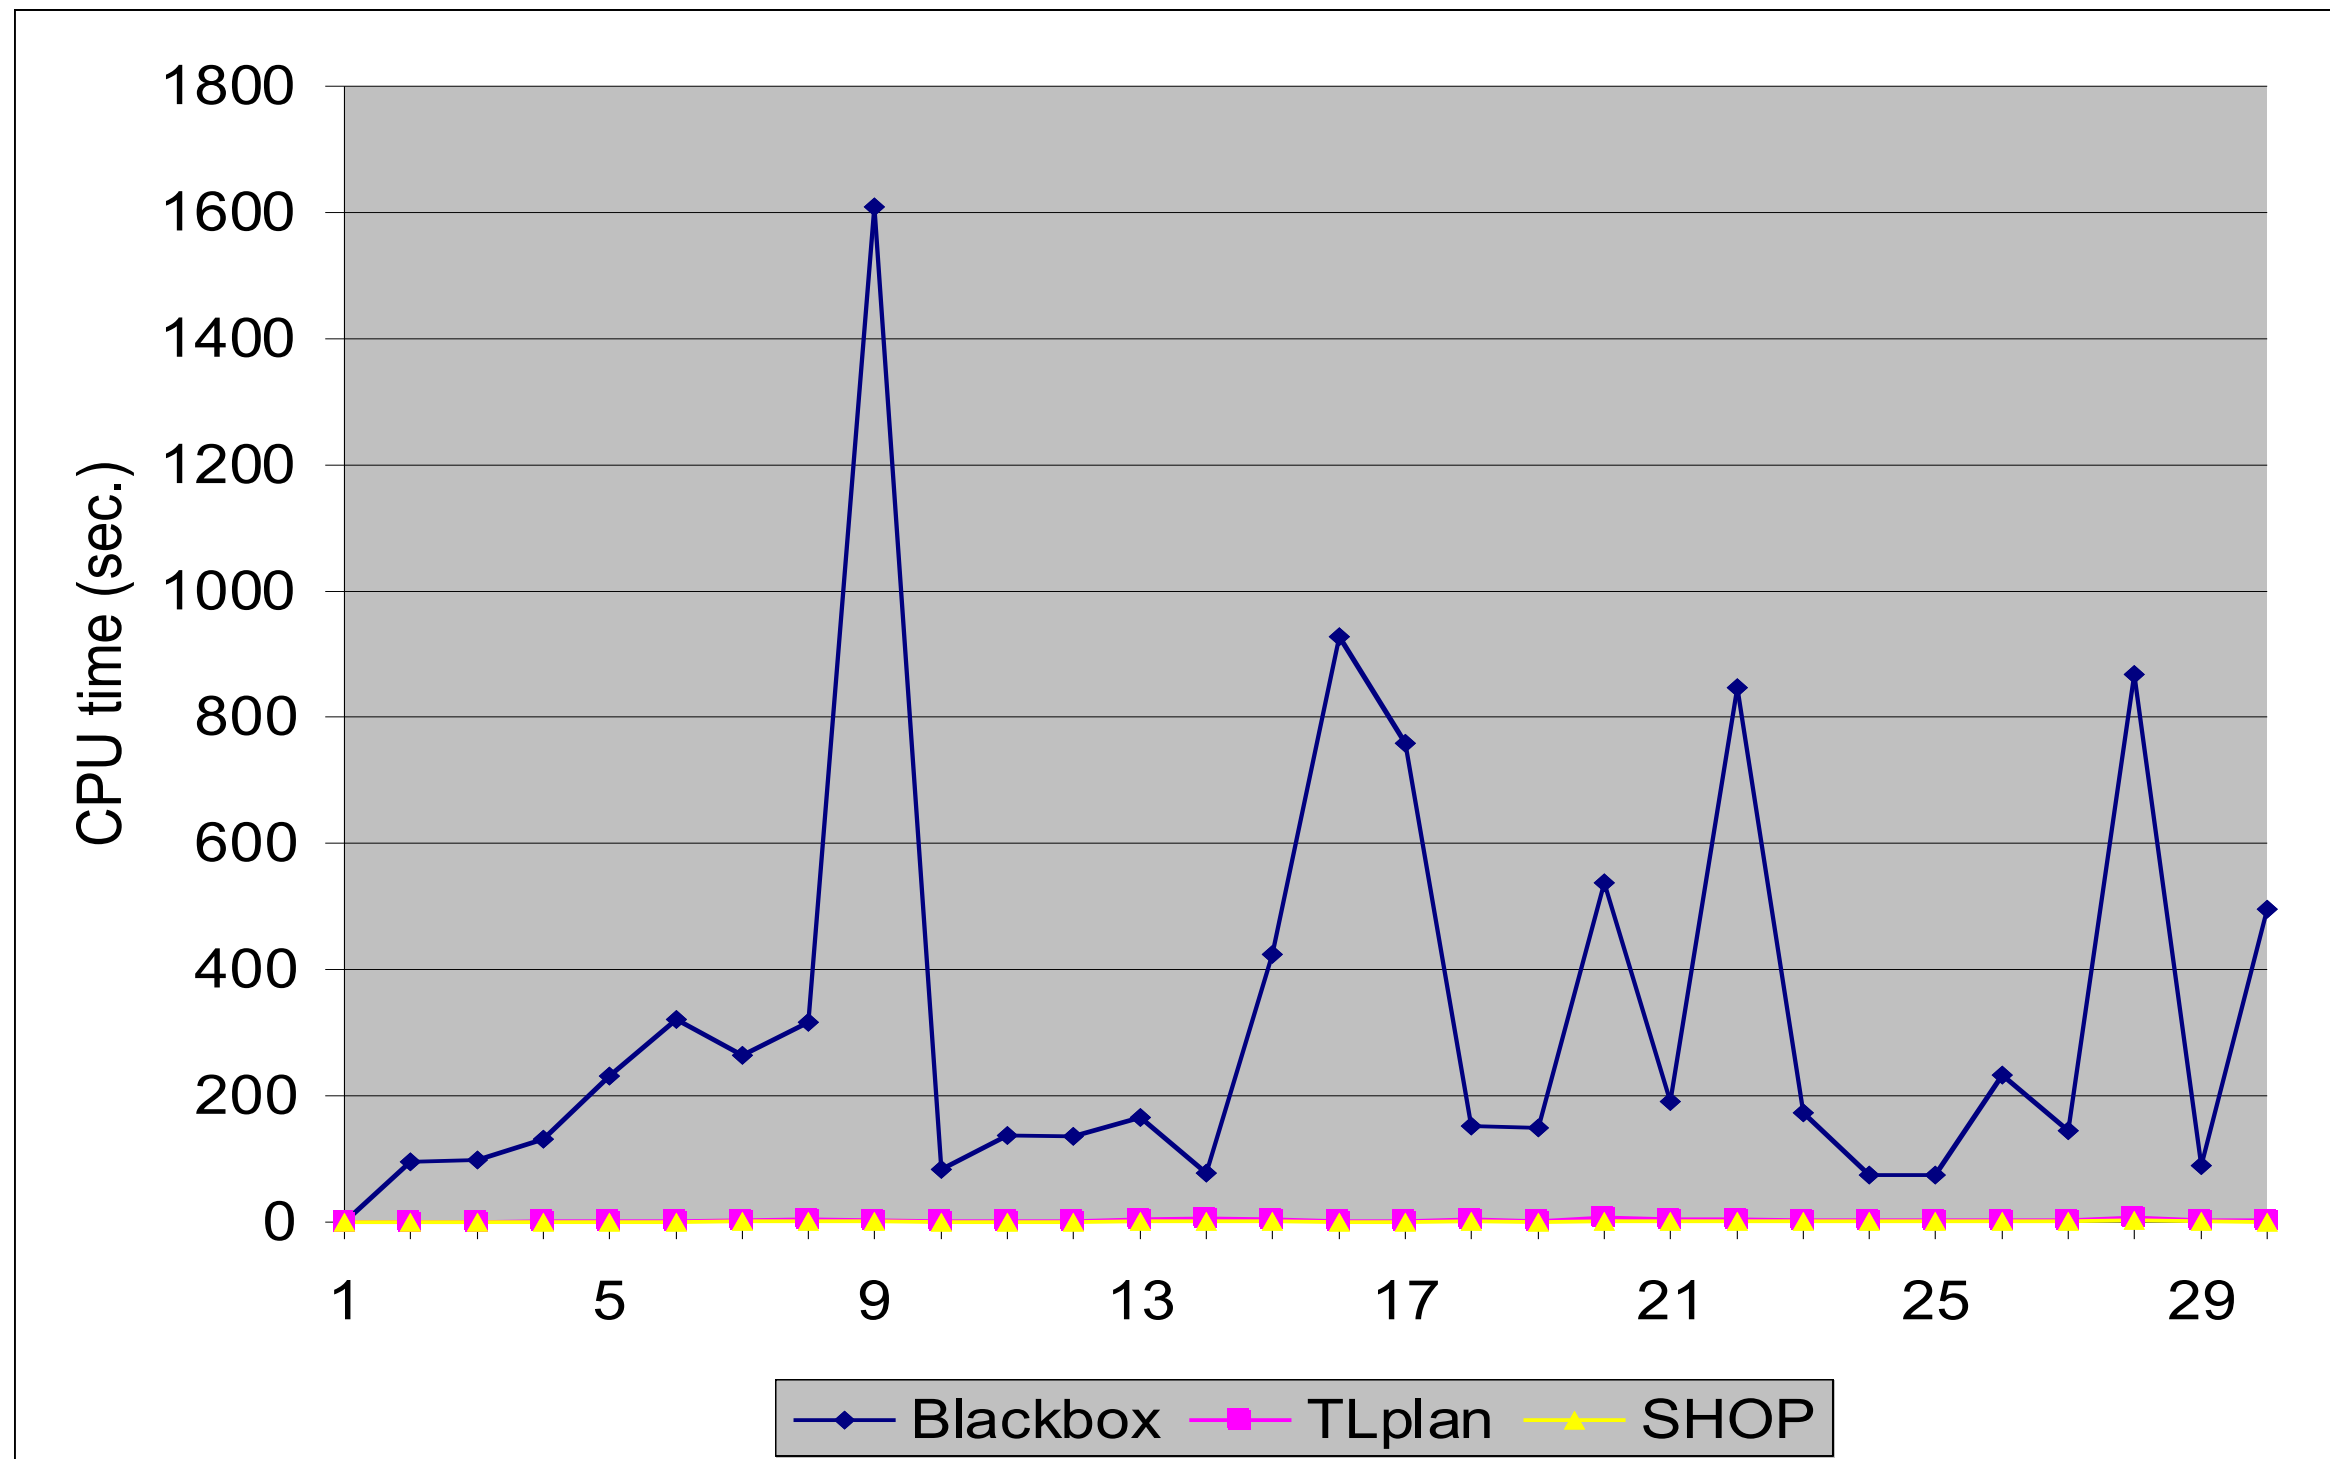
- Equivalent expressive power
- Both know the current state at each step of the planning process, and use this to prune operators
- Both can call external subroutines
  - SHOP uses “eval” to call LISP functions
  - In TLPlan, a function symbol can correspond to a computed function
- Main difference
  - in SHOP and SHOP2, the methods talk about what can be done
    - SHOP and SHOP2 don't do anything unless a method says to do it
  - TLPlan's control rules talk about what cannot be done
    - TLPlan will try everything that the control rules don't prohibit
- Which approach is more convenient depends on the problem domain



# Experimental Comparison


- Several years ago, we did a comparison of SHOP, TLPlan, and Blackbox
  - Blackbox is a domain-independent planner that uses a combination of Graphplan and satisfiability
  - One of the two fastest planners in the 1998 planning competition
- Test domain: the logistics domain
  - A classical planning problem
    - Much simpler than real logistics planning
  - Scenario: use trucks and airplanes to deliver packages
  - Like a simplified version of the DWR domain in which containers don't get stacked on each other
- Test conditions
  - SHOP and TLPlan on a 167-MHz Sun Ultra with 64 MB of RAM
  - We couldn't run Blackbox on our machine
  - Published results: Blackbox on a faster machine with 8 GB of RAM

# Logistics Domain Results




# Logistics Domain Results (continued)

- Same graph as before, but on a logarithmic scale




Average	Blackbox	TLPlan	SHOP
CPU time	327.1	2.9	1.1



Average no. of actions	Blackbox	TLPlan	SHOP
	82.5	54.5	51.9

# Summary: Results

- TLPlan and SHOP took similar amounts of time
  - In this experiment, SHOP was slightly faster, but in others TLPlan may be faster
- Blackbox took about 1000 times as much time and needed about 100 times as much memory
- Reasons why:
  - SHOP's input included domain-specific methods & axioms
  - TLPlan's input included domain-specific control rules
    - This enabled them to find near-optimal solutions in low-order polynomial time and space
  - Blackbox is a fully automated planner
    - No domain-specific knowledge
    - trial-and-error search, exponential time and space



# Domain-Configurable Planners Compared to Classical Planners

## ■ Disadvantage:

- Writing a knowledge base can be more complicated than just writing classical operators

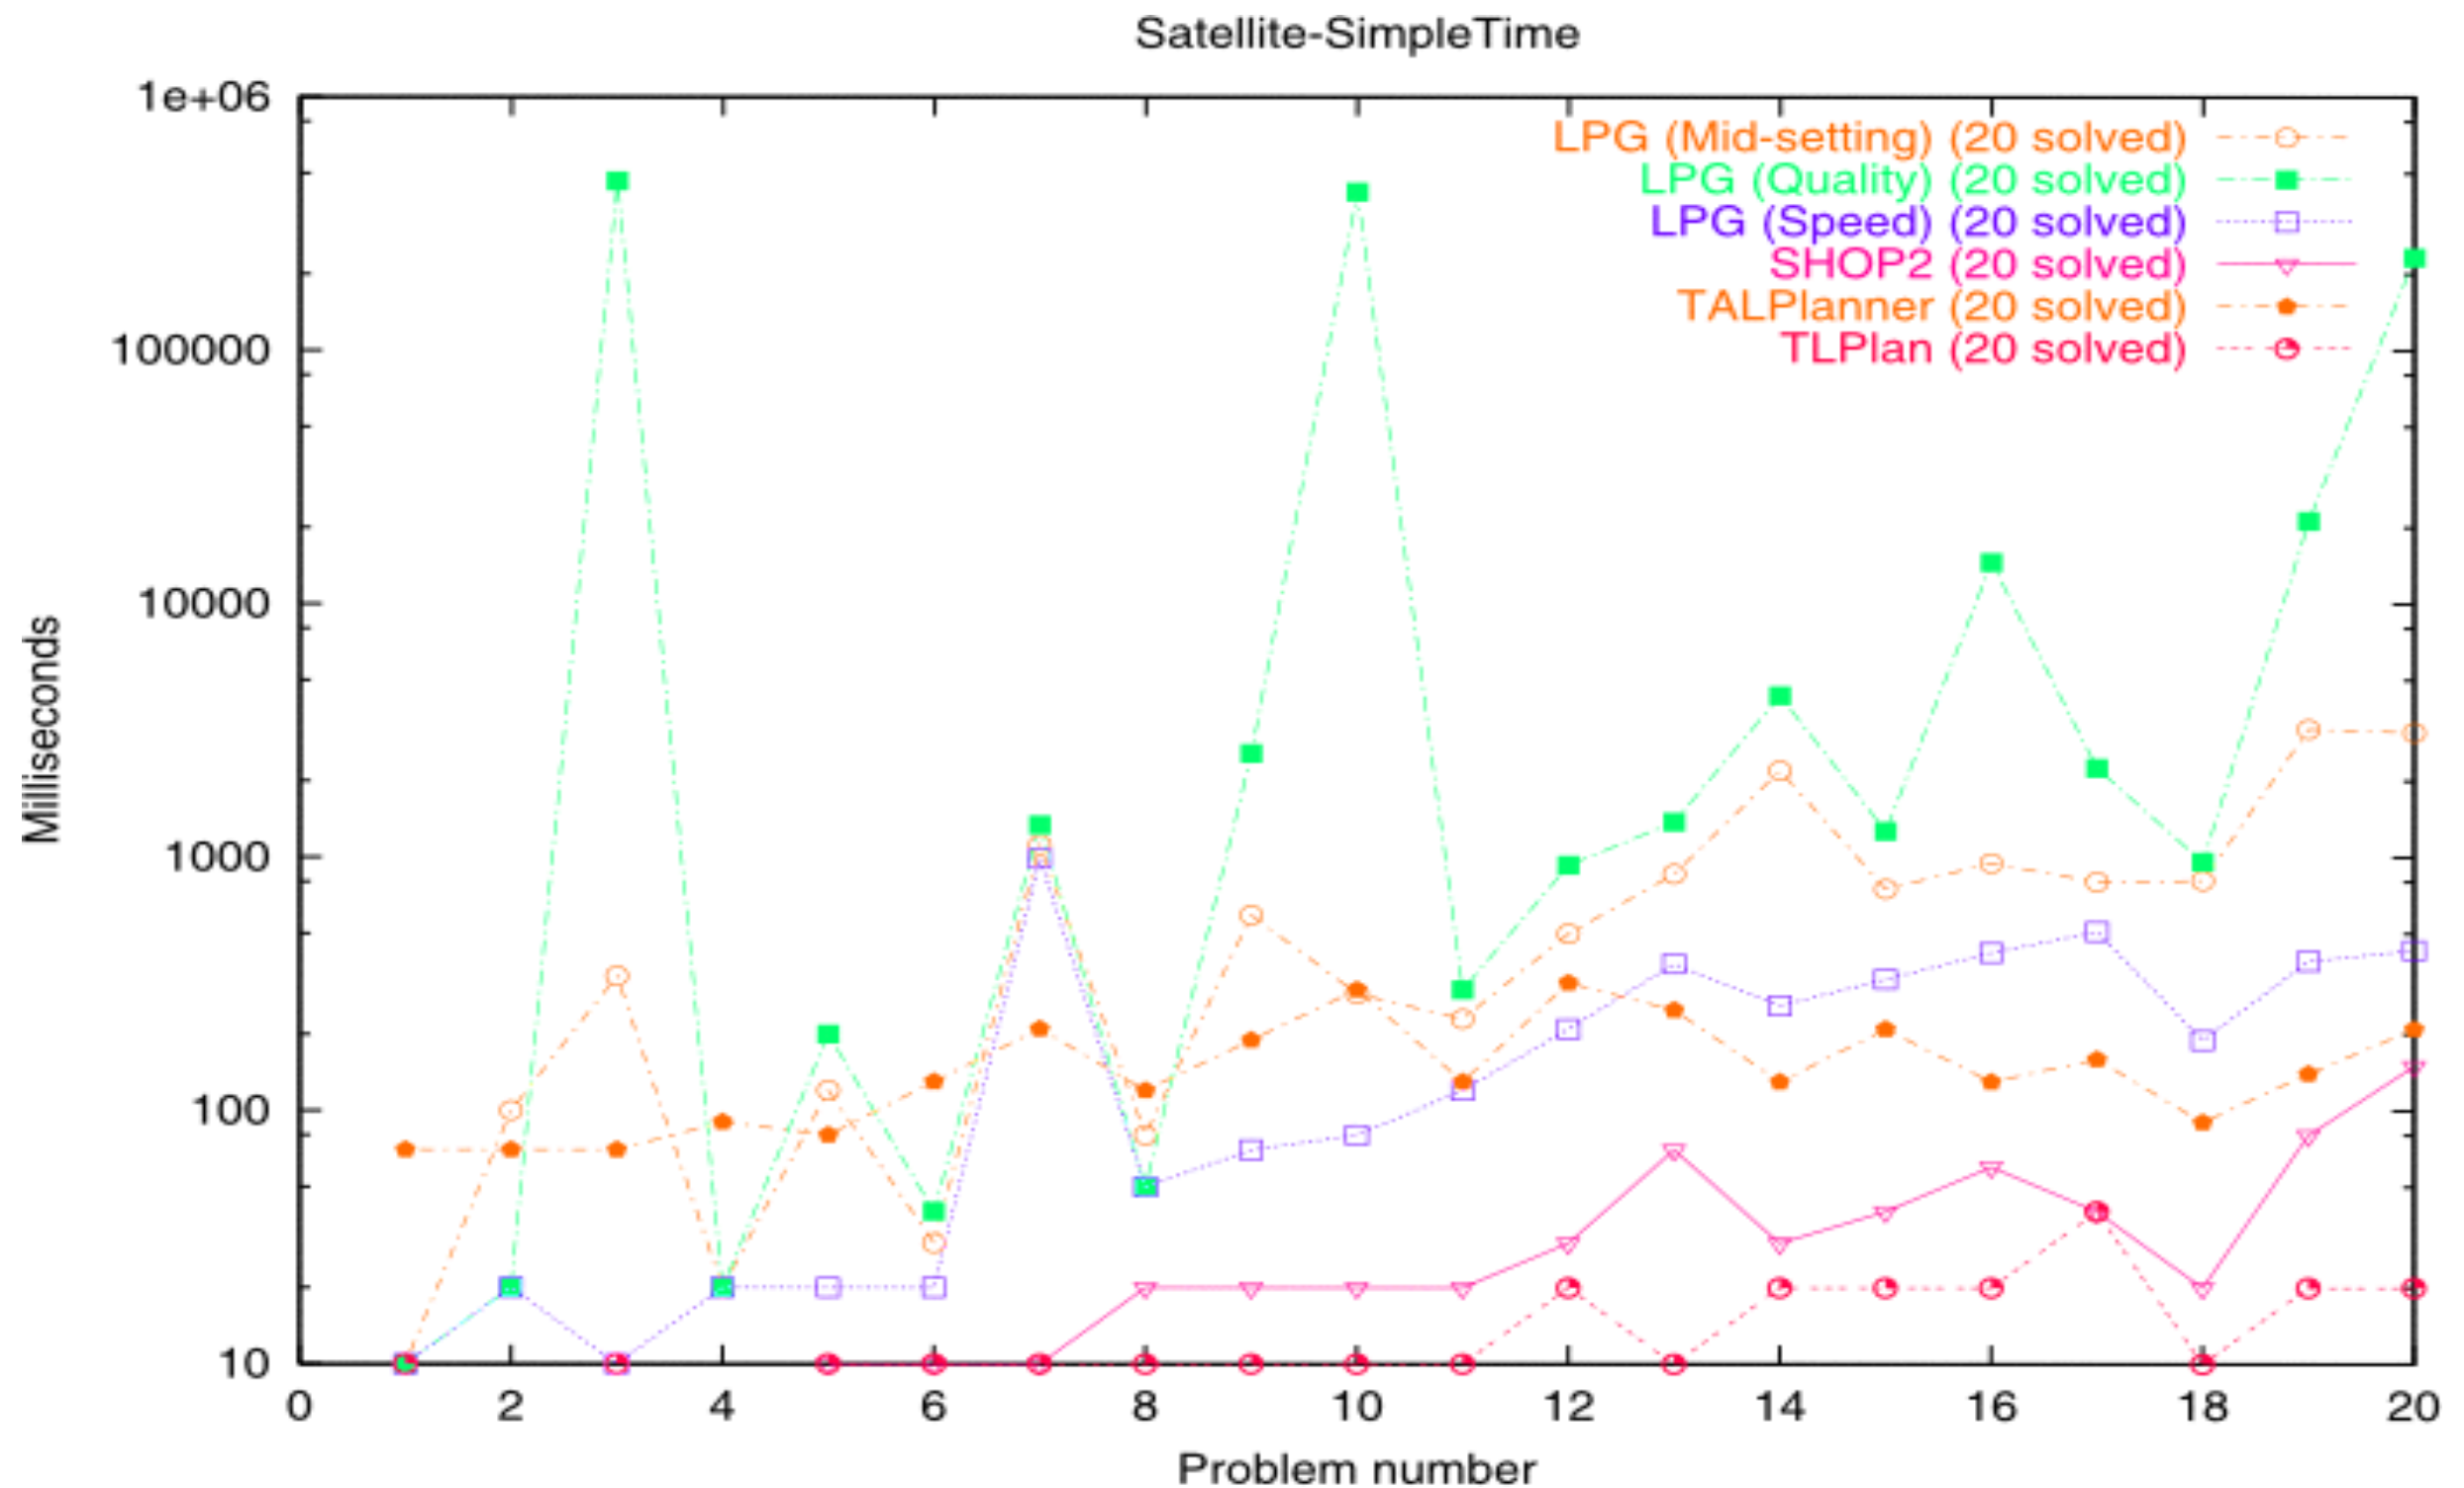
## ■ Advantage:

- We can encode “recipes” as collections of methods and operators
  - Express things that can’t be expressed in classical planning
  - Specify standard ways of solving problems
    - Otherwise, the planning system would have to derive these again and again from “first principles,” every time it solves a problem
  - Can speed up planning by many orders of magnitude (e.g., polynomial time versus exponential time)

# Example from the AIPS-2002 Competition


- The satellite domain
  - Planning and scheduling observation tasks among multiple satellites
  - Each satellite equipped in slightly different ways
- Several different versions. Results are shown for the following:
  - Simple time:
    - concurrent use of different satellites
    - data can be acquired more quickly if they are used efficiently
  - Numeric:
    - fuel costs for satellites to slew between targets; finite amount of fuel available.
    - data takes up space in a finite capacity data store
    - Plans are expected to acquire all the necessary data at minimum fuel cost.
  - Hard Numeric:
    - no logical goals at all – thus even the null plan is a solution
    - Plans that acquire more data are better – thus the null plan has no value
    - None of the classical planners could handle this

# Satellite Problem Domain: Simple Time: Runtime







# Satellite Problem Domain: Simple Time: Plan Quality




# Satellite Problem Domain: Numeric: Runtime



# Satellite Problem Domain: Numeric: Plan Quality



# Satellite Problem Domain: Hard Numeric: Runtime



# Satellite Problem Domain: Hard Numeric: Plan Quality

